

F5 Daily Briefing - Tuesday, May 12, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

21

ANNOUNCEMENTS


12

ACTIVE


5

ACTIVE CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (21)

PRODUCT LAUNCH

PALO ALTO NETWORKS

2026-05-12

Palo Alto Networks Launches Idira Identity Security Platform

Palo Alto Networks unveiled Idira, a next-generation identity security platform built on technology acquired in the \$25B CyberArk deal. Idira extends privileged access management to humans, machines, and AI agents under a single control plane, applying zero standing privilege and just-in-time access by default. Built-in AI surfaces hidden entitlements, flags unmanaged accounts, and recommends least-privilege remediation. Positioned alongside Strata and Cortex as the company's third core platform, Idira is generally available now with additional capabilities planned later in 2026. The launch responds to a reported 109:1 ratio of machine to human identities and the rapid rise of autonomous agents in production environments.

Source: <https://www.paloaltonetworks.com/blog/2026/05/idira-journey-democratize-privilege-controls/>

PRODUCT UPDATE

PALO ALTO NETWORKS

2026-05-12

Palo Alto Networks Expands Secure AI Agents with MCP Data Access Controls

Palo Alto Networks announced expanded controls and visibility for Model Context Protocol (MCP) data access within its Secure AI Agents product. The update tackles security blind spots created when AI agents use MCP servers without consistent authentication, governance, or visibility. New capabilities map which agent acted on which resource with associated permissions and targets, and add governance for MCP servers to mitigate prompt injection, supply-chain compromise, credential theft, over-scoped tokens, and remote code execution via compromised MCPs. The release strengthens Palo Alto's Prisma AIRS AI security stack against emerging agentic AI threats in enterprise deployments.

Source: <https://www.paloaltonetworks.com/blog/identity-security/secure-ai-agents-controls-visibility-mcp-data-access/>

PRODUCT LAUNCH**CISCO**

2026-05-12

Cisco Unveils Security Innovations for Enterprise AI Agent Adoption

Cisco introduced end-to-end security for agentic AI at RSA Conference 2026, organized around three pillars: protect the world from agents, protect agents from the world, and detect and respond to AI incidents at machine speed. Zero Trust Access is extended to AI agents with new Duo IAM capabilities, MCP policy enforcement, and intent-aware monitoring in Cisco Secure Access. Agents are registered in Duo IAM and mapped to accountable human owners for traceability, while Cisco Identity Intelligence provides visibility into agentic and non-human entities. The release targets the trust and governance barriers slowing enterprise agent deployment.

Source: <https://web-release.com/cisco-introduces-security-innovations-to-support-enterprise-ai-agent-adoption/>

PRODUCT LAUNCH**CISCO**

2026-05-12

Cisco Open-Sources Foundry Security Spec for Agentic Security Evaluation

Cisco published the Foundry Security Spec, an open-source, model- and stack-agnostic blueprint for building agentic security evaluation systems. The spec defines eight core agent roles (Orchestrator, Detector, Validator, Reporter, and others), a findings lifecycle, coordination substrate, and roughly 130 functional requirements, paired with a constitution of eleven inviolable principles. It is designed to work with GitHub's spec-kit for spec-driven development and complements Cisco's open-source CodeGuard project. The goal is to move security from noisy alerts to verifiable, auditable findings with operator-defined coverage and economic-yield thresholds, giving CISOs a defensible framework for evaluating LLM-driven security agents.

Source: <https://blogs.cisco.com/ai/announcing-foundry-security-spec>

CORPORATE NEWS**CLOUDFLARE**

2026-05-12

Cloudflare Cuts 20% of Workforce in AI-Driven Restructuring

Cloudflare announced layoffs of roughly 1,100 employees (about 20% of staff) as part of an AI-driven restructuring framed not as cost-cutting but as a reshape for the 'agentic AI era.' CEO Matthew Prince cited a 600%+ surge in internal AI usage over three months, with 97% of R&D staff using AI coding tools and autonomous agents reviewing code. Departing employees receive severance through end of 2026, US healthcare through year-end, and accelerated equity vesting. The cuts span nearly all functions except quota-carrying sales. Cloudflare beat Q1 2026 revenue and earnings expectations (\$639.8M, +34% YoY), but shares fell over 20% on softer near-term guidance.

Source: <https://www.securityweek.com/cloudflare-lays-off-1100-employees-in-ai-driven-restructuring/>

PARTNERSHIP**CROWDSTRIKE**

2026-05-12

CrowdStrike Extends Cloud Detection and Response to Google Cloud

CrowdStrike expanded its Cloud Detection and Response (CDR) capabilities to Google Cloud and extended the Falcon platform to regional Google Cloud infrastructure. The integration enables AI-native cloud security with unified CNAPP coverage, eliminating log batch processing delays in favor of real-time runtime detection. Google Cloud executives highlighted improved speed and precision in stopping cloud attacks, with telemetry flowing directly into Falcon for behavioral analytics and automated response. The move tightens CrowdStrike's multi-cloud posture and positions Falcon to compete more directly with cloud-native security tools as enterprises scale Google Cloud AI workloads.

Source: <https://www.techpartner.news/news/crowdstrike-adds-cdr-capabilities-to-google-cloud-625790>

SECURITY ADVISORY**ZSCALER**

2026-05-12

Zscaler Issues Protection for 12 New May 2026 Windows Vulnerabilities

Zscaler published a security advisory detailing protections deployed for 12 newly disclosed Windows vulnerabilities (CVE-2026-33835 through CVE-2026-41103) covered in Microsoft's May 2026 Patch Tuesday. The flaws span elevation-of-privilege and remote-code-execution issues across the kernel, Win32k, TCP/IP, WinSock, Cloud Files Mini Filter Driver, RDS, CDFS, and Word, affecting Windows 10/11 and Server 2019–2025 across x64 and ARM64. Through the Microsoft MAPP program, Zscaler proactively pushed Advanced Cloud Sandbox and Advanced Threat Protection signatures and will monitor for in-the-wild exploitation.

Source: <https://www.zscaler.com/br/security-advisories/zscaler-protects-against-12-new-vulnerabilities-windows-may-2026>

PRODUCT UPDATE**ZSCALER**

2026-05-12

Zscaler Implements OpenAI GPT-5.5-Cyber Under TAC Program

Zscaler announced it has stepped up to OpenAI's GPT-5.5-Cyber model under the OpenAI Trusted Access Cybersecurity (TAC) program, having previously participated with GPT-5.4-Cyber. The frontier model is being applied to AI Red Teaming for hardening customer AI deployments, generating optimized fixes when weaknesses are detected, and accelerating detection and remediation across Zscaler's daily transaction volume. The post frames defender strategy around two requirements: a Zero Trust architecture that makes applications invisible to attackers, plus frontier AI working on the defender's side to identify and fix issues at machine speed.

Source: <https://www.zscaler.com/blogs/company-news/next-advancement-openai-zscaler-implements-gpt-5-5-cyber>

PRODUCT LAUNCH**FORTINET**

2026-05-12

Fortinet Advances Security Operations Platform with Unified SOC and Agentic AI

At Accelerate 2026, Fortinet previewed FortiSOC, a cloud-delivered SOC service integrating FortiAnalyzer, FortiSIEM, FortiSOAR, and FortiTIP into a single platform with a unified data model, behavioral analytics, case management, and identity-focused investigations. FortiAI is expanded across the SecOps stack to power agentic workflows for automated alert triage, investigation, threat hunting, and response. FortiEndpoint gains enhancements, and FortiGuard managed services expand with broader third-party log sources and deeper Security Fabric integration. The goal is to unify telemetry, analytics, and response under one architecture with simplified licensing and elastic cloud scaling.

Source: <https://www.fortinet.com/corporate/about-us/newsroom/press-releases/2026/fortinet-advances-its-security-operations-platform-with-unified-soc-agentic-ai-and-expanded-endpoint-security>

SECURITY ADVISORY**FORTINET**

2026-05-12

Fortinet Patches Critical RCE Flaws in FortiSandbox and FortiAuthenticator

Fortinet issued security updates for two critical vulnerabilities. CVE-2026-44277 in FortiAuthenticator is an improper-access-control flaw allowing unauthenticated attackers to execute arbitrary code via crafted requests; it is patched in versions 6.5.7, 6.6.9, and 8.0.3 (Cloud is unaffected). CVE-2026-26083 in FortiSandbox is a missing-authorization flaw enabling unauthenticated RCE on FortiSandbox, FortiSandbox Cloud, and FortiSandbox PaaS web UIs. Fortinet says no in-the-wild exploitation has been observed, but urges rapid patching given the company's history of being targeted by ransomware and espionage actors.

Source: <https://www.bleepingcomputer.com/news/security/fortinet-warns-of-critical-rce-flaws-in-fortisandbox-and-fortiauthenticator/>

PARTNERSHIP

FORTINET

2026-05-12

Fortinet Deepens FortiAI Gate Integration with NVIDIA to Secure Enterprise AI

Fortinet announced a deeper integration of FortiAI Gate with NVIDIA to secure AI workloads and autonomous agents across data centers and cloud environments. The joint solution provides zero-trust security, real-time governance, and low-latency protection for LLMs and AI agents while supporting data sovereignty through self-hosted deployments on NVIDIA GPUs and Dynamo. Highlighted benefits include accelerated and safe AI runtimes, on-prem AI sovereignty, and governance that reduces threat exposure and response time. The integration positions FortiAI Gate squarely against AI gateway and runtime-security players as enterprises scale agentic deployments.

Source: <https://www.globenewswire.com/news-release/2026/05/12/3292915/0/en/Fortinet-Deepens-Integration-to-Uniquely-Secure-Enterprise-AI-at-Scale-with-NVIDIA.html>

PRODUCT UPDATE

AWS

2026-05-12

AWS Security Agent Adds Full Repository Code Review

AWS added a full-repository code-review capability to AWS Security Agent, performing deep, context-aware analysis across an entire codebase to find systemic vulnerabilities that pattern-matching tools miss. When issues are detected, the agent delivers precise remediation tied to exact file and line numbers, accelerating fix workflows. The feature is free during preview for existing AWS Security Agent customers and is available in all regions the agent supports. It is enabled from the AWS Security Agent console with quickstart guidance in documentation, extending AWS's developer-side security tooling into broader application-security territory.

Source: <https://aws.amazon.com/about-aws/whats-new/2026/05/aws-security-agent-full-repository-code-review/>

PRODUCT LAUNCH

AWS

2026-05-12

Claude Platform on AWS Reaches General Availability

AWS announced general availability of Claude Platform on AWS, giving customers direct access to Anthropic's full Claude Platform—APIs, console, and beta features—through their existing AWS accounts. The platform is operated by Anthropic with data processed outside the AWS security boundary, suitable for teams without strict regional residency requirements. No separate Anthropic accounts are needed: authentication uses AWS IAM, billing flows through AWS Marketplace consumption, and activity is auditable via CloudTrail. Features include Claude Managed Agents, advisor strategy, web search/fetch, code execution, files API, Agent Skills, MCP connector, prompt caching, citations, and batch processing.

Source: <https://aws.amazon.com/blogs/aws/introducing-claude-platform-on-aws-anthropics-native-platform-through-your-aws-account/>

PRODUCT LAUNCH

AWS

2026-05-12

Amazon Redshift Introduces Graviton-Based RG Instances with Integrated Data Lake Query Engine

AWS launched a new Redshift RG instance family powered by AWS Graviton, delivering up to 2.2x faster performance than RA3 at roughly 30% lower price per vCPU. The instances ship with an integrated data-lake query engine, enabling SQL analytics across both the data warehouse and S3-based data lakes with up to 2.4x faster Apache Iceberg and 1.5x faster Parquet queries. The family is tuned for high query volumes and low-latency analytics, including AI agent workloads. Example mappings: ra3.xlplus maps to rg.xlarge (4 vCPUs, 32 GB) and ra3.4xlarge maps to rg.4xlarge.

Source: <https://aws.amazon.com/blogs/aws/amazon-redshift-introduces-aws-graviton-based-rg-instances-with-an-integrated-data-lake-query-engine/>

PRODUCT LAUNCH**HPE / JUNIPER NETWORKS**

2026-05-12

HPE Unifies Private Cloud and Data Platforms for AI Data Readiness

HPE expanded its GreenLake hybrid cloud stack with HPE Private Cloud Gen4, adding Kubernetes management alongside VMs in a single orchestration layer with an upgrade path to HPE Morpheus for hybrid and multi-cloud needs. The release ships on ProLiant Compute Gen12 servers and adds HPE Alletra Storage MP X10000, which unifies file and object storage—scaling to 16 nodes and 23 PB raw, with RDMA-accelerated file storage to speed AI pipelines. New agentic AI management capabilities land across Alletra MP B10000, Zerto Software, and Data Fabric Software, with Veeam and Zerto integrations for backup and recovery to ease VMware migrations.

Source: <https://www.hpe.com/us/en/newsroom/press-release/2026/05/hpe-delivers-unified-private-clouds-and-data-platforms-to-accelerate-enterprise-modernization-and-ai-data-readiness.html>

PRODUCT LAUNCH**CITRIX (NETSCALER)**

2026-05-12

Citrix Introduces Platform Flex for Persona-Based Secure Access Consumption

Citrix announced Citrix Platform Flex, a secure access platform that moves from fixed licensing to a flexible, credit-based consumption model. Flex credits are allocated against workforce personas—developers, traders, knowledge workers, and others—to tailor performance, security, and resilience per worker type. The platform bundles managed desktops, enterprise browser, and zero-trust access, with support for hybrid and multi-cloud deployments and AI-ready workloads. Citrix positions it as a way to cut overprovisioning, align spend with actual usage, and adapt faster to changing workforce demands without added complexity.

Source: <https://www.helpnetsecurity.com/2026/05/12/citrix-moves-secure-access-to-a-flexible-credit-based-consumption-model/>

SECURITY ADVISORY**CITRIX (NETSCALER)**

2026-05-12

NetScaler Issues Critical Patches for CVE-2026-3055 and CVE-2026-4368

Citrix released a security bulletin for two NetScaler ADC and NetScaler Gateway vulnerabilities. CVE-2026-3055 (CVSS 9.3) is an input-validation flaw causing memory overread when the appliance is configured as a SAML IdP. CVE-2026-4368 (CVSS 7.7) is a race condition that mixes up user sessions on configurations using AAA Vservers or Gateway/VPN Vservers. Fixed builds include NetScaler 14.1-60.58, 14.1-66.59, 13.1-62.23, and 13.1-FIPS/NDcPP 13.1.37.262. Citrix urges immediate upgrades and configuration audits to identify whether SAML IdP profiles or affected vservers are in use.

Source: <https://support.citrix.com/external/article/CTX696300>

EARNINGS**FASTLY**

2026-05-12

Fastly Q1 2026 Revenue Up 20% on Security and AI Edge Demand

Fastly reported record Q1 2026 results with revenue of \$173 million, up 20% year over year, beating estimates. Security revenue jumped 47% to \$38.8 million (22% of total), driven by new API discovery, bot management, and DDoS offerings. Compute and 'other' revenue rose 67% to \$8 million on AI workloads. Non-GAAP gross margin reached a record 65.1%, with \$19.1 million of non-GAAP operating income and \$22.9 million net profit. Net retention climbed to 113% and RPO hit \$369 million (+63% YoY). Management raised full-year revenue guidance to \$710–\$725 million, though shares fell sharply on softer near-term outlook.

Source: <https://www.harianbasis.co/en/fastly-exceeds-q1-2026-earnings-expectations>

THREAT RESEARCH

HIDDENLAYER

2026-05-12

HiddenLayer Uncovers Malicious Hugging Face Repo Posing as OpenAI Release

HiddenLayer disclosed a malicious Hugging Face repository (Open-OSS/privacy-filter) that typosquatted an OpenAI Privacy Filter release and pushed Windows infostealer malware. The fake repo cloned OpenAI's release page and included a loader.py that disabled SSL verification, fetched payload instructions via jsonkeeper.com, ran PowerShell, and downloaded a Rust-based stealer with a Microsoft-Edge-update persistence trick. Targets included Chromium/Firefox browser data, Discord storage, crypto wallets, FileZilla configs, and Telegram sessions. The repo logged ~244,000 downloads and trended highly within 18 hours. HiddenLayer found six other repos using nearly identical loader logic and warned that AI model registries are now a supply-chain attack vector.

Source: <https://www.infosecurity-magazine.com/news/malicious-hugging-face-repo/>

THREAT RESEARCH

BARRACUDA NETWORKS

2026-05-12

Barracuda Releases 2026 Email Threats Report on AI and Phishing-as-a-Service

Barracuda published its 2026 Email Threats Report, finding that AI-driven social engineering and phishing-as-a-service are driving both volume and effectiveness of email attacks. Key data: 1 in 3 emails are malicious or spam; 48% of malicious mail is phishing; 34% of companies suffer at least one account takeover per month; over 10% of HTML attachments are malicious; 70% of malicious PDFs hide QR codes leading to phishing sites; and 90% of high-volume phishing campaigns use PhaaS kits. Attackers are shifting from file-based payloads to URL-based delivery, using QR codes in trusted documents and AITM techniques to bypass defenses.

Source: <https://blog.barracuda.com/2026/05/12/ai-phishing-service-changing-email-threat-landscape>

THREAT RESEARCH

BARRACUDA NETWORKS

2026-05-12

Barracuda's Mythos Hype Index Questions AI's Real Impact on Vulnerability Discovery

Barracuda introduced the Mythos Hype Index, a daily measure that compares 2026 CVE growth against historical baselines to test whether AI-assisted vulnerability discovery is materializing. Early findings: 2026 CVE growth is only slightly above historical trend, and fewer than 200 CVEs explicitly credit AI/LLM tooling. The post argues that disclosure-pipeline constraints and under-attribution help explain the gap, but that practical risk management—patch prioritization, asset visibility, and exposure reduction—remains a stronger predictor of real-world risk than raw CVE volume. The piece pushes back on AI-vulnerability-discovery hype for IT and security leaders.

Source: <https://blog.barracuda.com/2026/05/12/mythos-hype-index-ai-vulnerability-discovery>