

F5 Daily Briefing — Thursday, May 14, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS
TRACKED

14

ANNOUNCEMENTS


9

ACTIVE


4

ACTIVE
CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (14)

PRODUCT LAUNCH

AWS

2026-05-14

AWS Announces General Availability of EC2 M3 Ultra Mac Instances

AWS released general availability of Amazon EC2 M3 Ultra Mac instances, built on Apple M3 Ultra Mac Studio hardware and powered by the AWS Nitro System. Targeted at developers running heavy build/test workloads for iOS, macOS, iPadOS, tvOS, watchOS, visionOS, and Safari, the instances feature 28-core CPUs, 60-core GPUs, 32-core Neural Engines, and 256GB unified memory. Versus M4 Max instances, M3 Ultra offers ~2x unified memory, ~1.75x CPU cores, ~1.5x GPU cores, and ~2x Neural Engine cores, plus up to 10 Gbps networking and 8 Gbps EBS. Available in US East (N. Virginia) and US West (Oregon).

Source: <https://aws.amazon.com/about-aws/whats-new/2026/05/amazon-ec2-m3-ultra-mac-instances-generally-available/>

PRODUCT LAUNCH**AWS**

2026-05-14

Amazon Bedrock Introduces Advanced Prompt Optimization and Migration Tool

AWS rolled out an Advanced Prompt Optimization and Migration tool for Amazon Bedrock, designed to help customers migrate prompts to new models or boost performance on current ones. The tool runs a metric-driven feedback loop: users supply a prompt template, example inputs, ground-truth answers, and an evaluation metric, and Bedrock iteratively optimizes the prompt. Outputs include the original and final prompt templates, evaluation scores, cost estimates, and latency measurements. The feature targets prompt-driven tasks like document analysis and large-scale inference workflows, with sample code available on GitHub to help teams accelerate model migration and AI delivery optimization.

Source: <https://aws.amazon.com/blogs/aws/amazon-bedrock-introduces-new-advanced-prompt-optimization-and-migration-tool/>

ACQUISITION**Akamai**

2026-05-14

Akamai to Acquire LayerX for ~\$205M to Expand Zero Trust with AI Usage Control

Akamai announced its intent to acquire LayerX, a provider of browser-based AI usage control and secure enterprise browser technology, for approximately \$205 million. The deal extends Akamai's Zero Trust security portfolio into the browser, giving customers real-time visibility and governance over how employees interact with generative AI, SaaS AI tools, and autonomous agents — without requiring a proprietary browser or major infrastructure changes. LayerX's team joins Akamai's Zero Trust organization in Tel Aviv. The transaction is expected to close in Q3 2026, dilute non-GAAP EPS by ~\$0.12 in fiscal 2026, and contribute roughly \$10M in ARR by year-end.

Source: <https://www.akamai.com/newsroom/press-release/akamai-technologies-announces-intent-to-acquire-layerx-advancing-its-workforce-security-strategy-with-ai-usage-control>

PRODUCT LAUNCH**Barracuda Networks**

2026-05-14

Barracuda Brings Automated Threat Response to the Email Inbox

Barracuda Networks introduced new automated response capabilities for its email security suite, designed to speed and streamline how inbound threats are handled directly in user inboxes. The feature integrates into Barracuda's broader email defense product line and aims to reduce manual triage workload for security teams, improve mean time to respond, and contain account-takeover and phishing scenarios faster. The announcement details rollout plans, automated reply use cases, and configuration steps customers can take within the Barracuda platform. It complements Barracuda's recent 2026 Email Threats Report findings on AI-driven phishing and phishing-as-a-service driving the need for layered, automated defenses.

Source: <https://blog.barracuda.com/2026/05/14/barracuda-email-automated-threat-response>

EARNINGS**Cisco**

2026-05-14

Cisco Beats Q3 Estimates and Cuts ~4,000 Jobs to Fund AI Investment

Cisco reported fiscal Q3 results that beat expectations: adjusted EPS of \$1.06 (vs. \$1.04 expected) and revenue of \$15.84B (vs. \$15.56B expected), up 12% YoY. Networking jumped 25% to \$8.82B while security stayed roughly flat near \$2B. Cisco raised its full-year AI revenue target to \$4B (from \$3B) and AI-related orders guidance to \$9B (from \$5B), with \$5.3B in AI hyperscaler orders YTD. Cisco also announced cutting fewer than 4,000 roles (under 5% of workforce) and ~\$1B in pre-tax restructuring charges to reallocate spending toward AI and cybersecurity. Shares rose ~15% after hours.

Source: <https://www.cnn.com/2026/05/13/cisco-csco-q3-earnings-report-2026.html>

PRODUCT LAUNCH**Cisco**

2026-05-14

Cisco Open-Sources Foundry Security Spec for Agentic AI Evaluation

Cisco open-sourced its Foundry Security Spec, a model- and stack-agnostic blueprint for building agentic security evaluation systems. The framework, derived from Cisco's internal security engineering work, is designed to help defenders move from noisy alerts to verifiable security findings and scale vulnerability identification at machine speed. The spec combines agent-based harnesses and orchestration with human expertise, exposing roles such as Indexer, Cartographer, Detector, Triager, Validator, Coverage-Guide, and Reporter. It works with GitHub spec-kit workflows and is publicly available on CiscoDevNet/foundry, with Cisco inviting practitioners to build and contribute their own evaluation systems.

Source: <https://blogs.cisco.com/ai/announcing-foundry-security-spec>

REPORT**CrowdStrike**

2026-05-14

CrowdStrike Releases 2026 Financial Services Threat Landscape Report

CrowdStrike released its 2026 Financial Services Threat Landscape Report, detailing how DPRK-nexus adversaries stole billions in digital assets in 2025 and are industrializing cybercrime with AI-powered deception. Hands-on-keyboard intrusions against financial institutions rose 43% globally and 48% in North America over the past two years, with adversaries increasingly exploiting trusted identities and SaaS apps to evade legacy defenses. The report cites a 51% YoY rise in digital asset theft and 423 financial firms named on leak sites (up 27%). CrowdStrike urges intelligence-led visibility, continuous hunting, and AI-enabled defense via Counter Adversary Operations on the Falcon platform.

Source: <https://ir.crowdstrike.com/news-releases/news-release-details/crowdstrike-2026-financial-services-threat-landscape-report>

EARNINGS**Fortinet**

2026-05-14

Fortinet Reports Strong Q1 2026 Results; Raises Full-Year Guidance

Fortinet reported Q1 2026 revenue up 20% YoY to \$1.85B, with product revenue surging 41% to \$645M and billings up 31% to \$2.09B — all above guidance. GAAP operating margin reached 31% (non-GAAP 36%), and free cash flow hit \$1.01B. The company raised full-year 2026 revenue guidance to roughly \$7.71B-\$7.87B (~15% growth) and lifted non-GAAP EPS guidance to \$3.10-\$3.16. Fortinet also touted its seventh straight Gartner Peer Insights SD-WAN Customers' Choice recognition, reinforcing platform momentum across SASE, cloud, and AI-driven security demand.

Source: <https://www.zawya.com/en/press-release/companies-news/fortinet-reports-strong-first-quarter-2026-financial-results-gwgcizih>

PRODUCT LAUNCH**Fortinet**

2026-05-14

Fortinet Introduces FortiOS 8.0 with Secure AI Controls, Flexible SASE, and Simplified SD-WAN

At Fortinet Accelerate 2026, Fortinet unveiled FortiOS 8.0, the next major release of its flagship Secure Networking operating system. The update introduces secure AI controls, Fabric-Based AI Agents, flexible SASE deployment, and a simplified SD-WAN architecture designed to reduce complexity and improve operational efficiency. FortiOS 8.0 also brings quantum-safe capabilities and broader AI adoption support, positioning Fortinet's platform to scale protection across hybrid digital infrastructure. New hardware including FortiGate 3500G and FortiGate 400G expands the platform footprint at the high end and at the branch.

Source: <https://www.fortinet.com/corporate/about-us/newsroom/press-releases/2026/fortinet-introduces-fortios-8-expand-secure-networking-with-secure-ai-controls-fabric-based-ai-agents-flexible-sase-and-simplified-sdwan>

PRODUCT LAUNCH**Fortinet**

2026-05-14

Fortinet Advances Security Operations Platform with Unified SOC, Agentic AI, and Expanded Endpoint Security

Fortinet announced major updates to its Security Operations Platform at Accelerate 2026, including a preview of FortiSOC, a cloud-delivered service unifying FortiAnalyzer, FortiSIEM, FortiSOAR, and FortiTIP. Expanded FortiAI capabilities deliver agentic workflows that automate alert triage, investigation, threat hunting, and cross-tool response over a unified data model. FortiEndpoint now offers single-agent unification across ZTNA, SASE, EPP, EDR, and DLP with AI-based application visibility. Fortinet's SOC architecture also adds managed detection and response options via FortiGuard, with simplified licensing and elastic cloud scaling for self-managed, cloud, and managed deployments.

Source: <https://www.fortinet.com/corporate/about-us/newsroom/press-releases/2026/fortinet-advances-its-security-operations-platform-with-unified-soc-agentic-ai-and-expanded-endpoint-security>

PARTNERSHIP**Fortinet**

2026-05-14

Fortinet Enhances FortiAI Gate with NVIDIA-Accelerated AI Security Capabilities

Fortinet integrated its FortiAI Gate platform with NVIDIA AI infrastructure to secure enterprise AI deployments at scale. The combined solution delivers zero-trust security for AI workloads, data, and autonomous agents across cloud and data center environments, with real-time visibility and governance over AI activity. FortiAI Gate uses inline, GPU-accelerated security on NVIDIA Blackwell GPUs with the Dynamo framework, targeting high throughput with low latency. It supports self-hosted deployments for AI sovereignty and leverages NVIDIA Nemotron safety models to monitor prompts, defend against prompt injection and data leakage, and enforce zero-trust guardrails across multi-tenant AI environments.

Source: <https://industrialcyber.co/news/fortinet-enhances-fortiaigate-platform-with-nvidia-accelerated-ai-security-capabilities/>

PARTNERSHIP**HPE / Juniper Networks**

2026-05-14

HPE Taps TD Synnex and Ingram Micro as Global Distribution Backbone in Post-Juniper Channel Reset

Hewlett Packard Enterprise selected TD Synnex and Ingram Micro as its global distribution partners following the acquisition of Juniper Networks, marking a major channel reset for the combined HPE Networking portfolio. The move aims to streamline the post-merger channel model, enable full-portfolio selling across networking, cloud, and AI, and improve partner support, training, and consistency. HPE will retain regional and specialist distributors while elevating these two as the global backbone. The plan includes expanding distributor expertise across the entire HPE stack and creating a single integrated distribution landscape designed to drive growth and unified go-to-market for HPE's expanded networking offering.

Source: <https://www.crn.com/news/channel-news/2026/hpe-taps-td-synnex-ingram-micro-as-global-distribution-backbone-in-post-juniper-channel-reset>

CORPORATE**Radware**

2026-05-14

Radware Announces New \$80 Million Share Repurchase Plan

Radware's board authorized a new \$80 million share buyback program, the 2026 Plan, which expires March 15, 2027. Management may repurchase ordinary shares in open-market or privately negotiated transactions subject to market conditions and compliance with U.S. and Israeli law, including Rule 10b-18 and potentially Rule 10b5-1 plans. The authorization does not obligate Radware to repurchase any specific number of shares, and the program may be suspended or terminated at management's discretion. The plan reflects continued capital return to shareholders alongside Radware's investments in its AI-driven application security and DDoS protection portfolio.

Source: <http://www.radware.com/newsevents/pressreleases/2026/radware-announces-new-80-million-share-repurchase-plan/>

Zscaler: "While You Embrace AI, Fix This First" — Zero Trust Foundations for AI Era

Zscaler published guidance warning that organizations must shore up exposure and lateral movement risks before scaling AI adoption, or AI will amplify existing vulnerabilities rather than capability. The post argues that exposed resources — public IPs, open ports, reachable services — are continuously profiled and exploitable at machine speed, including AI models themselves. AI agents can autonomously traverse endpoints and data sources, turning a local compromise into systemic risk. Zscaler frames its Zero Trust Exchange as the architectural answer: verify everything, trust nothing by default, and treat reachability control plus zero-trust segmentation as prerequisites to enterprise AI rollout.

Source: <https://www.zscaler.com/blogs/product-insights/while-you-embrace-ai-fix-this-first>