

F5 Daily Briefing — Monday, May 18, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

9

ANNOUNCEMENTS


7

ACTIVE


4

ACTIVE CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (9)

PRODUCT LAUNCH

FORTINET

2026-05-16

Fortinet Unveils FortiSOC with Agentic AI at Accelerate 2026

At Fortinet Accelerate 2026 on May 16, Fortinet announced FortiSOC, a cloud-delivered unified SOC service that fuses FortiAnalyzer, FortiSIEM, FortiSOAR, and FortiTIP into a single console with a shared data model. FortiAI agentic workflows now extend across the SecOps stack, automating alert triage, investigation, and threat hunting via Model Context Protocol for continuity. FortiEndpoint also adds a single-agent unification across ZTNA, SASE, EPP, EDR, and DLP, positioning Fortinet to counter AI-driven threats.

Source: <https://www.fortinet.com/corporate/about-us/newsroom/press-releases/2026/fortinet-advances-its-security-operations-platform-with-unified-soc-agentic-ai-and-expanded-endpoint-security>

PRODUCT LAUNCH

PALO ALTO NETWORKS

2026-05-15

Idira Identity Security Platform for the AI Enterprise Picks Up Industry Coverage

Computer Weekly and additional outlets profiled Palo Alto Networks' newly launched Idira identity security platform over the weekend, framing it as a unified privileged access management (PAM) control plane for human, machine, and agentic AI identities. Idira eliminates standing privileges by default, applies AI-driven least-privilege recommendations, and integrates with Prisma AIRS 3.0 for AI agent security. The platform extends Palo Alto's portfolio reach to CyberArk customers and consolidates a fragmented identity market behind a single AI-era control plane.

Source: <https://www.computerweekly.com/blog/CW-Developer-Network/Palo-Alto-Networks-Idira-A-next-gen-platform-for-human-agentic-machine-identity-security>

PRODUCT LAUNCH

CLOUDFLARE

2026-05-15

Cloudflare Launches Workflows V2 with Deterministic Execution and 50K Concurrent Workflows

On May 15, Cloudflare released Workflows V2, introducing a deterministic, replayable execution model for stateful multi-step workflows including AI inference pipelines. The platform scales to 50,000 concurrent workflow instances (up from 4,500) and 300 new executions per second, with step-level tracing, fan-out/fan-in parallelism, and durable AI agent submissions. Built on Workers, Queues, and Durable Objects, V2 targets long-running AI/agent and data-sync pipelines where idempotent, recoverable execution is essential.

Source: <https://www.infoq.com/news/2026/05/cloudflare-workflows-v2-release/>

SECURITY UPDATE

CLOUDFLARE

2026-05-15

Cloudflare WAF Ships Emergency Rule for nginx CVE-2026-42945

Cloudflare's WAF deployed an emergency managed rule on May 15 to detect exploitation attempts against nginx CVE-2026-42945 in the rewrite module. The release joined a broader May 12-16 changelog that expanded R2 SQL with JOIN/subquery support, refreshed Cloudflare Access login UX, added NAT-T support for IKE on UDP 500 in Magic Transit, and rolled out a refreshed Gateway natural-language firewall-policy builder, underscoring the breadth of Cloudflare's application-security platform cadence.

Source: <https://releasebot.io/updates/cloudflare>

FINANCIAL

CISCO

2026-05-15

Cisco CEO Frames Company as 'Critical Infrastructure for the AI Era' After Earnings Beat

In a May 15 Fortune interview, CEO Chuck Robbins positioned Cisco as the backbone of AI-era infrastructure, citing a 13%+ share rally and record stock price following the latest earnings beat. Robbins emphasized that the multi-year pivot toward AI networking, security, and platform consolidation is finally paying off, with hyperscaler AI demand pulling Cisco gear into more deployments. The piece reinforces Cisco's broader platformization narrative across Splunk, Secure Access, and AI Defense.

Source: <https://fortune.com/2026/05/15/critical-infrastructure-ai-era-cisco-ceo-earnings-beat-shares-record/>

STRATEGIC

ARISTA NETWORKS

2026-05-15

Arista Doubles Down on AI Fabric, Targeting \$3.5B AI Revenue in 2026

Coverage on May 15 detailed Arista's two-front strategy to diversify its customer base while standardizing AI-era data-center hardware. Arista guided to roughly \$11.5B in 2026 revenue with AI fabric revenue doubling to about \$3.5B, joined a multi-source agreement with AMD, Cisco, and Meta on Expanded Beam Optics (EBO) for AI clusters, and flagged Oracle as a potential ~10% revenue customer. A May 16 follow-up highlighted the 7800-based universal AI spine and liquid-cooled XPO optics as key AI-networking accelerants.

Source: <https://www.ad-hoc-news.de/boerse/news/ueberblick/arista-networks-stages-a-two-front-offensive-diversifying-customers-and/69346117>

ACQUISITION

AKAMAI

2026-05-15

Akamai's \$205M LayerX Acquisition Spotlights Browser-Based AI Usage Control

Industry coverage from May 15-17 expanded on Akamai's \$205M agreement to acquire LayerX, framing the deal as a major Zero Trust pivot into browser-native AI usage governance. LayerX monitors and controls employee interactions with generative AI, SaaS AI services, and autonomous agents inside mainstream browsers without forcing a proprietary enterprise browser. Reports highlighted Akamai's fourth Tel Aviv acquisition in five years; close is expected Q3 2026 with roughly \$10M ARR contribution and ~\$0.12 dilution to FY26 non-GAAP EPS.

Source: <https://www.leadrpro.com/blog/akamai-layerx-deal-puts-ai-sales-tools-under-browser-scrutiny>

PRODUCT LAUNCH

AWS

2026-05-15

AWS Pushes Graviton Into the Redshift Analytics Stack With New RG Instances

On May 15, AWS launched Graviton-powered RG instances for Amazon Redshift, integrating data warehouse and data lake analytics into a single query engine. AWS reports the new RG family delivers up to 2.2x faster warehouse performance than RA3, roughly 30% lower cost per vCPU, and gains of up to 2.4x on Apache Iceberg and 1.5x on Parquet workloads. The launch signals AWS's expansion of custom silicon further up the AI analytics stack, supporting massive inference and feature-engineering pipelines.

Source: <https://www.datacenterknowledge.com/infrastructure/aws-pushes-graviton-into-redshift-analytics-stack>

PRODUCT LAUNCH

AWS

2026-05-15

AWS Adds Advanced Prompt Optimization to Amazon Bedrock

AWS rolled out Advanced Prompt Optimization (APO) for Amazon Bedrock on May 15, giving teams a managed tool to evaluate prompts against user-defined datasets across up to five foundation models and surface the highest-performing configurations per workload. APO is generally available across multiple regions with enterprise pricing tied to underlying service costs. The addition strengthens Bedrock's AI inference delivery story by closing the gap between model selection, prompt design, and production performance for agentic AI applications.

Source: <https://www.infoworld.com/article/4171870/aws-adds-advanced-prompt-optimization-tool-to-bedrock.html>