

F5 Daily Briefing – Monday, July 20, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

8

ANNOUNCEMENTS


5

ACTIVE


4

ACTIVE CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (8)

PARTNERSHIP

Palo Alto Networks

2026-07-17

Palo Alto Networks and AT&T Launch Quantum-Resilient SASE Fabric

Palo Alto Networks and AT&T Business unveiled a joint Quantum-Resilient SASE Fabric that integrates Prisma SD-WAN with AT&T's global network to deliver quantum-ready secure connectivity. It embeds native post-quantum cryptography into Prisma SASE with a crypto-agile, software-defined design, dual classical/quantum-resistant identities, and zero-touch provisioning. A Service Provider Interconnect links Palo Alto's security platform to AT&T's Private Network Core to cut latency, targeting NIS2, DORA and CNSA 2.0 compliance without a forklift upgrade.

<https://www.paloaltonetworks.com/blog/2026/07/palo-alto-networks-and-att-delivering-quantum-resilient-sase-fabric/>

PARTNERSHIP

Palo Alto Networks

2026-07-17

IBM, Red Hat and Palo Alto Networks Expand Project Lightwell

Palo Alto Networks joined IBM and Red Hat to expand Project Lightwell, accelerating vulnerability response across open source, commercial apps, OT and connected devices. Palo Alto adds network-level Virtual Patching to complement Lightwell's open-source remediation, enabling same-day virtual patch deployment to shorten the window from discovery to protection. The collaboration includes shared vulnerability intelligence, secure information-sharing, anonymized exploit telemetry, and IBM Security Services advisory support to prioritize and validate protections across complex environments.

<https://mid-east.info/ibm-red-hat-and-palo-alto-networks-expand-project-lightwell-to-help-organizations-respond-to-software-vulnerabilities/>

PRODUCT LAUNCH

Palo Alto Networks

2026-07-17

Prisma AIRS AI Gateway Reaches General Availability as Enterprise AI Control Plane

Palo Alto Networks announced general availability of its Prisma AIRS AI Gateway, now integrating Portkey technology to serve as a unified enterprise AI control plane sitting inline between AI interactions, model providers and agentic apps. It unifies execution across LLM, MCP and A2A gateways with claimed sub-millisecond routing, 99.999% availability and 68 trillion tokens processed monthly. The GA delivers governance, identity and runtime controls at scale, positioning Prisma AIRS as the foundation for secure AI infrastructure against rivals Zscaler and CrowdStrike.

<https://vendordeep.com/report/palo-alto-networks-launches-gateway-centralized-control>

REPORT

Check Point Software

2026-07-17

Check Point Research Publishes AI Security Report 2026

Check Point Research released its AI Security Report 2026, warning that AI has shifted from assisting attackers to autonomously operating live intrusions. It cites a breach in which two AI tools compromised nine Mexican government agencies by executing 5,317 AI-driven commands across 34 sessions, and notes the vulnerability-to-exploit window has shrunk from days to hours. Adverse AI prompts doubled year over year; roughly 87-93% of organizations face at least one high-risk AI interaction monthly. Check Point urges Security for AI, Security by AI, and Security with AI.

<https://research.checkpoint.com/2026/ai-security-report-2026/>

SECURITY UPDATE

Cloudflare

2026-07-17

Cloudflare Ships Emergency WAF Release Blocking WordPress RCE and SQLi

Cloudflare issued an emergency Web Application Firewall release blocking active exploitation of two high-severity WordPress vulnerabilities: an unauthenticated remote code execution via the REST API and a related SQL injection. New managed rules (Unauthenticated RCE and SQLi, set to Block) were deployed at 17:03 UTC to all customers, including free and paid plans. WordPress released fixes in 7.0.2 with backports; Cloudflare stressed WAF coverage is not a substitute for patching and advised reviewing rules and monitoring Security Events.

<https://developers.cloudflare.com/changelog/post/2026-07-17-emergency-waf-release/>

PRODUCT UPDATE

Cloudflare

2026-07-18

Cloudflare Advances 'Pay Per Use' Model for AI Answers

Cloudflare detailed its Pay Per Use model that compensates publishers when their content surfaces in AI-generated answers rather than per page fetch. Starting September 15, 2026, Cloudflare will block many ad-supported, mixed-use AI crawlers by default unless site owners opt in, pushing AI companies to pay for content used in training and answers. Early partners include Ceramic.ai and You.com. CEO Matthew Prince framed the

shift as essential to sustaining publishing economics amid AI-driven discovery and heavily non-human web traffic.

<https://interviewer99.com/cloudflare-pay-per-use-ai-crawlers-publishers/>

PRODUCT LAUNCH

Akamai

2026-07-17

Akamai Deploys NVIDIA 'AI Grid' Across 4,400 Edge Locations

Akamai said it deployed NVIDIA's AI Grid reference design across its 4,400 edge locations, claiming to be first to operationalize the offering at scale. The rollout uses thousands of NVIDIA RTX Pro 6000 Blackwell Server Edition GPUs at edge sites to enable low-latency, cost-efficient real-time AI inference. Highlighted use cases include real-time video, physical AI, agentic and personalized experiences, and access to fine-tuned models. Akamai argues edge inference offers the right balance of latency, cost and scale for long-tail AI workloads versus centralized AI factories.

<https://www.datacenterdynamics.com/en/news/akamai-deploys-nvidia-ai-grid-across-4400-edge-locations-claims-to-be-first/>

PRODUCT UPDATE

AWS

2026-07-17

SageMaker HyperPod Adds Partition-Level Topology for Slurm Clusters

AWS announced that Amazon SageMaker HyperPod now supports partition-level network topology for Slurm-orchestrated clusters, letting a single cluster mix block topology (e.g., EC2 UltraServer ml.p6e-gb200.36xlarge) and tree topology (ml.p5 instances) partitions. HyperPod auto-optimizes topology per partition and adapts during scale operations, with topology-aware scheduling on by default. It requires Slurm 25.11+ and supported GPU instances, and is available in all Regions where HyperPod runs. The feature improves distributed training performance, GPU-to-GPU communication and NCCL efficiency for large AI workloads.

<https://aws.amazon.com/about-aws/whats-new/2026/07/hyperpod-partition-topology-slurm/>