

F5 Daily Briefing - Wednesday, August 05, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

22


ANNOUNCEMENTS

12

ACTIVE

5

ACTIVE CATEGORIES



ALL ANNOUNCEMENTS (22)

PRODUCT LAUNCH

PALO ALTO NETWORKS

2026-08-05

Cortex Cloud Adds Software Supply Chain Security Module

Palo Alto Networks introduced a Software Supply Chain Security module inside Cortex Cloud to build trust into the SDLC. It adds Software Supply Chain Trust Scores, a continuous integrity metric evaluating SBOM dependencies and the build environment, plus a Supply Chain Attack Threat Center tracking emerging supply chain threats and CVEs in real time. The module extends Cortex Cloud's code-to-cloud coverage, reinforcing Palo Alto's platformization strategy of folding adjacent security functions into a single suite.

<https://www.paloaltonetworks.com/blog/cloud-security/software-supply-chain-security-module-new/>

RESEARCH

PALO ALTO NETWORKS

2026-08-05

Palo Alto Networks at Black Hat: How AI Erased the 50-Day Patch Window

At Black Hat 2026 Palo Alto Networks argued AI-driven vulnerability discovery has collapsed the traditional patch window. Unit 42's agentic AI research system NOVA surfaced roughly 14,090 previously unreported vulnerabilities across 3,915 open-source projects in about two months. The company pitched PAN-OS 12.2 Ceres and its Advanced Virtual Patching capability as the answer, shielding unpatched systems at the network layer while fixes are developed and deployed.

<https://www.networkworld.com/article/4205156/palo-alto-networks-at-black-hat-how-ai-erased-the-50-day-patch-window.html>

PRODUCT LAUNCH

PALO ALTO NETWORKS

2026-08-05

Prisma AIRS: Unified Data Protection for Claude via Inference Hooks

Palo Alto Networks announced the Prisma AIRS API now integrates directly with Claude Enterprise through Anthropic's inference hooks, classifying prompts in real time before data is accessed. Existing enterprise DLP policies, confidentiality definitions and compliance dashboards automatically govern Claude usage with no new rules to author, and Claude events flow into the same audit reporting as other channels. The company positions Claude as one more governed channel, tying the news to Prisma AIRS AI Gateway general availability and Black Hat 2026 demos.

<https://www.paloaltonetworks.com/blog/2026/08/prisma-air-unified-data-protection-for-claude/>

SECURITY ADVISORY**CISCO**

2026-08-05

Cisco IOS XE Software Security Hardening Release: August 2026

Cisco published a security hardening release for IOS XE Software addressing multiple internally discovered vulnerabilities found during a comprehensive internal security review. Issues are grouped by CWE class with a single CVE per grouping, and Cisco reported no evidence of active exploitation. No workarounds are available, so customers must upgrade to fixed releases (for example 17.9 to 17.9.10). The advisory was part of Cisco's August 5 PSIRT bundle covering IOS, IOS XE, IMC, RoomOS and Terminal Services Agent.

<https://sec.cloudapps.cisco.com/security/center/content/CiscoSecurityAdvisory/cisco-sa-hardening-iosxe-V8NMuMZJ>

SECURITY ADVISORY**CISCO**

2026-08-05

Cisco Catalyst SD-WAN Software Security Hardening Release: August 2026

Cisco issued a companion hardening release for Catalyst SD-WAN Software as part of its August 5 PSIRT advisory bundle. The release remediates multiple internally discovered vulnerabilities identified through Cisco's internal security review process, grouped by CWE class. Cisco directs customers to upgrade to the listed fixed software releases, noting there are no workarounds. The bundled disclosure approach spans Cisco's networking and collaboration portfolio, underscoring the maintenance burden that accompanies broad platform consolidation.

<https://sec.cloudapps.cisco.com/security/center/content/CiscoSecurityAdvisory/cisco-sa-hardening-sdwan-faLcR3K>

OTHER**CROWDSTRIKE**

2026-08-05

CrowdStrike Announces \$100,000 International AI Security Challenge with AWS

CrowdStrike, with AWS, launched "AI Unlocked: Agents of Chaos," a \$100,000 international AI red-teaming challenge. The virtual, interactive game teaches defenders how prompt injection, agent hijacking and rogue AI behavior can be exploited and stopped in enterprise AI agents. It runs in three acts — Sanctum (\$10,000), Gatekeeper (\$20,000) and Basilisk (\$70,000) — from August 31 to September 29, 2026, with on-site play at Fal.Con 2026 in Las Vegas. The program showcases CrowdStrike's Falcon AI Detection and Response capabilities.

<https://www.crowdstrike.com/en-us/press-releases/crowdstrike-announces-international-ai-security-challenge/>

AWARD**ZSCALER**

2026-08-05

Zscaler Wins 3 CRN Tech Innovator Awards

Zscaler announced three 2026 CRN Tech Innovator Awards: Security/Secure Access Service Edge for the Zero Trust Exchange, Security/Network Security Enterprise for Zscaler Internet Access, and Edge Computing and IoT for Zscaler Cellular. The company positioned the recognition as validation of its AI security platform, enterprise network security capabilities and edge connectivity portfolio. Zscaler highlighted the awards as a partner enablement moment for conversations around secure access for users, apps, AI agents and devices across distributed environments.

<https://www.zscaler.com/blogs/partner/zscaler-wins-3-crn-tech-innovator-awards>

PARTNERSHIP**ZSCALER**

2026-08-05

Zscaler Integrates With Claude Inference Hooks to Scale AI While Addressing Risks

Zscaler announced a collaboration with Anthropic to enforce Zero Trust policy directly inside Claude Enterprise using Claude inference hooks. A single policy defined in Zscaler is enforced at the model layer across all Claude touchpoints, and AI Guard evaluates conversations in real time to prevent data exposure rather than relying on after-the-fact audits. Zscaler frames the integration as eliminating governance fragmentation by giving security teams unified visibility and enforcement across the AI attack surface.

<https://www.zscaler.com/blogs/product-insights/zscaler-integrates-claude-inference-hooks-scale-ai-while-addressing-risks>

PRODUCT LAUNCH**CLOUDFLARE**

2026-08-05

Cloudflare OS Is the First AI Workspace Built Around How Companies Actually Work

Cloudflare launched Cloudflare OS, an open-source AI workspace running on its global network that gives every employee a secure, browser-based environment connected to internal company systems. It combines an isolated agent runtime for code execution, a Zero Trust security and governance framework for internal data access, and a platform for personal apps anyone can build and share. Cloudflare OS is available now on GitHub, with a managed deployment option coming to the dashboard and partner-led implementations from Presidio, Happy Cog and others.

<https://www.cloudflare.com/press/press-releases/2026/cloudflare-os-is-the-first-ai-workspace-built-around-how-companies-actually-work/>

PRODUCT LAUNCH **CLOUDFLARE** 2026-08-05

Cloudflare Gives Companies Full Visibility to Audit and Analyze AI Use

Cloudflare introduced Identity-Aware AI Gateway and User Insights, giving IT and security teams real-time visibility into which employees or AI agents are calling AI services. Teams can set individual or team spending limits and automatically flag risky content without changing existing workflows. User Insights learns normal AI usage patterns for every user and automated system, then alerts when behavior deviates from the baseline. The features target AI cost control, data-leak risk and policy enforcement while preserving developer velocity.

<https://www.cloudflare.com/press/press-releases/2026/cloudflare-gives-companies-full-visibility-to-audit-and-analyze-ai-use/>

PRODUCT LAUNCH **AWS** 2026-08-05

Amazon DynamoDB Now Supports Real-Time Vector Search at Any Scale

AWS made native vector search generally available in Amazon DynamoDB, letting customers store embeddings alongside operational data without a separate vector store. It delivers single-digit millisecond latency with 99%+ recall, scales to trillions of vectors, supports up to 4096 dimensions with Euclidean, cosine and dot-product metrics, and allows attribute filtering at query time via a new SearchVectors API. Target use cases include semantic memory for AI agents, retrieval-augmented generation, product similarity and personalized recommendations, with serverless pay-as-you-use pricing.

<https://aws.amazon.com/blogs/aws/amazon-dynamodb-now-supports-real-time-vector-search-at-any-scale/>

PARTNERSHIP **AWS** 2026-08-05

AWS Partners with Anthropic and OpenAI to Bring AWS Continuum into Developer Workflows

AWS expanded AWS Continuum into coding environments through integrations with Anthropic Claude Code, OpenAI Codex and Kiro. Continuum for code vulnerabilities, now in preview, performs on-demand scanning, prioritizes findings using the customer's actual AWS context (IAM policies, configurations, network topology and exposure surfaces), validates exploitability in a sandbox, and returns prioritized remediation guidance to the coding assistant. An agent-team loop orchestrates model selection, environment access and secure code delivery so developers do not manage orchestration themselves.

<https://aws.amazon.com/blogs/security/aws-partners-with-anthropic-and-openai-to-bring-aws-continuum-into-developer-workflows/>

PRODUCT LAUNCH **AWS** 2026-08-05

AWS Reintroduces Network Firewall Proxy for Secure Egress Connectivity

AWS folded its explicit proxy capability directly into AWS Network Firewall rather than shipping it standalone, so one policy model and one firewall now cover both transparent filtering and explicit proxy inspection. A new no-source-preservation deployment mode attaches the firewall to a NAT Gateway, masking the client's original source IP and routing egress through the NAT address. This enables centralized egress control even with overlapping CIDRs and eliminates duplicate rule languages, unifying egress policy management under a single service.

<https://aws.amazon.com/blogs/networking-and-content-delivery/reintroducing-network-firewall-proxy-for-secure-egress-connectivity/>

RESEARCH **HIDDENLAYER** 2026-08-05

How To Secure Coding Agents: A Five-Layer Security Model for Agent Harnesses

HiddenLayer published research arguing the real attack surface for AI coding agents is the entire agent harness — prompts, tools, skills, MCP servers and orchestration logic — not just the model. It proposes a five-layer security model: visibility into every trust relationship, control to contain compromised components, validation instead of assumption, monitoring outside the model itself, and governance with clear ownership. The piece complements HiddenLayer's recently launched Agent Harness Security runtime product and positions whole-ecosystem defense as essential as enterprises scale AI-assisted development.

<https://www.hiddenlayer.com/research/how-to-secure-coding-agents>

PRODUCT LAUNCH **AKAMAI** 2026-08-05

Akamai Workforce Protector (Formerly LayerX) Advances Enterprise Security

Akamai launched Akamai Workforce Protector, the rebranded LayerX technology, as an interaction-layer security product governing how employees and AI agents use AI, SaaS, web and private applications. It delivers real-time data loss prevention, AI usage control, secure enterprise browser capabilities and visibility into modern desktop AI apps and agentic workflows, natively integrated with Akamai's global private access solution. Akamai positions it to close gaps left by traditional SSE, CASB and DLP tools and to stop sensitive data reaching unauthorized AI models.

<https://www.stocktitan.net/news/AKAM/akamai-workforce-protector-formerly-layer-x-advances-enterprise-et3pqjy3a39v.html>

RESEARCH

AKAMAI

2026-08-05

Akamai Research: Nearly Half of Enterprise AI Use Bypasses Corporate Security

Akamai released its 2026 State of the Internet security report, finding roughly 47% of enterprise AI interactions run through personal, unmanaged accounts, creating major shadow AI visibility gaps. The report details three AI-native attack techniques — Vibe hacking, Cursorjacking and Cometjacking — that bypass traditional perimeters by exploiting local environments, browser extensions and public web pages. It also notes nearly 75% of AI extensions request high or critical permissions, and recommends SSO-based shadow AI elimination, real-time prompt inspection and least-privilege controls for autonomous agents.

<https://www.stocktitan.net/news/AKAM/akamai-research-nearly-half-of-enterprise-ai-use-bypasses-corporate-hzghlxofnbj9.html>

FINANCIAL

FASTLY

2026-08-05

Fastly Announces Second Quarter 2026 Financial Results

Fastly reported Q2 2026 revenue of \$183.3 million, up 23% year over year, with record GAAP gross margin of 63.3% and non-GAAP gross margin of 65.8%. Security revenue grew 43% to \$41.7 million while Network Services rose 17% to \$133.9 million. LTM net revenue retention hit 117%, the highest in over three years, and RPO reached \$341 million, up 38%. Operating loss narrowed to \$14.4 million from \$36.9 million, and Fastly raised full-year revenue outlook to \$732-746 million.

<https://finviz.com/news/377939/fastly-announces-second-quarter-2026-financial-results>

RESEARCH

FASTLY

2026-08-05

Fastly Research: AI Traffic Grows 6.5x Faster Than Human Traffic

Alongside its Q2 results, Fastly published research showing AI and machine-generated traffic is growing 6.5 times faster than human traffic this year, arguing organizations need new strategies to manage machine-to-machine traffic at the edge. Fastly paired the finding with product and partnership momentum including a Skyfire collaboration for edge-enabled trusted commerce with real-time AI agent transactions, a new C++ SDK for Fastly Compute aimed at secure low-latency AI workloads, and an anti-piracy collaboration with LALIGA.

<https://www.stocktitan.net/news/FSLY/fastly-announces-second-quarter-2026-financial-8e1htq49ycp7.html>

FINANCIAL

A10 NETWORKS

2026-08-05

A10 Networks Reports Financial Results for the Second Quarter of 2026

A10 Networks reported Q2 2026 revenue of \$80.1 million, up 15.5% year over year, with first-half revenue of \$155.1 million. Non-GAAP gross margin was 80.3%, adjusted EBITDA margin 30.5%, and non-GAAP diluted EPS \$0.25. Management credited demand for security-focused, next-generation networking for AI infrastructure, the June 2026 completion of the TrojAI acquisition, and a multi-year Microsoft agreement supporting scaled AI deployments. Cash and marketable securities totaled \$357.3 million as of June 30, 2026.

<https://investor.wedbush.com/wedbush/article/bizwire-2026-8-5-a10-networks-reports-financial-results-for-the-second-quarter-of-2026>

FINANCIAL

A10 NETWORKS

2026-08-05

A10 Networks Raises Full-Year 2026 Outlook and Declares Quarterly Dividend

Alongside Q2 results, A10 Networks raised full-year 2026 guidance, lifting expected revenue growth to 12-14% year over year (from 10-12%) and EPS growth to 14-16% (from 12-14%). The company cited robust security and AI-infrastructure-related demand, the completed TrojAI acquisition, and an extended long-term Microsoft agreement. The board approved a quarterly dividend of \$0.06 per share, payable September 1, 2026 to stockholders of record on August 17, 2026.

<https://www.stocktitan.net/sec-filings/ATEN/8-k-a10-networks-inc-reports-material-event-f8dec2dea811.html>

FINANCIAL

ARISTA NETWORKS

2026-08-05

Arista Q2 Earnings Call Centers on Supply Gains and 2026 Guidance Raise

Arista's Q2 2026 earnings call highlighted the company's first \$3 billion quarter, with revenue of \$3.036 billion up 37.7% year over year on AI-driven data center networking demand. Executives pointed to easing supply constraints and raised full-year expectations, guiding Q3 revenue to roughly \$3.3 billion. Arista also emphasized its 1.6 Tbps AI fabric platforms, including liquid-cooled options, as central to scaling AI back-end networks.

<https://finance.yahoo.com/markets/stocks/articles/anet-q2-earnings-call-centers-140000921.html>

PRODUCT LAUNCH

IMPERVA (THALES)

2026-08-05

Thales Launches Imperva for AWS to Help Organizations Protect Applications and APIs

Thales announced Imperva for AWS, a SaaS-based Imperva Web Application Firewall deploying natively on Amazon CloudFront to protect web applications, APIs and client-side code without traffic leaving the AWS network. The offering adds bot and API protection, centralized policy management and visibility across AWS and multi-cloud environments, and claims zero added latency. It is available on AWS Marketplace with pay-as-you-go pricing, targeting faster onboarding and lower operational overhead for highly regulated enterprise workloads.

https://www.imperva.com/company/press_releases/thales-launches-imperva-for-aws-to-help-organizations-protect-applications-and-apis/