

F5 Daily Briefing - Monday, August 17, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

20

ANNOUNCEMENTS


15

ACTIVE


5

ACTIVE CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (20)

CUSTOMER WIN Zscaler 2026-08-17

Komatsu Australia Moves 1,800 Field Devices to Zscaler Cellular

Komatsu Australia is shifting roughly 1,800 field devices off Telstra-locked SIMs to Zscaler Cellular as part of a broader zero trust program built on Zscaler's Zero Trust Exchange. The network-agnostic MVNO approach lets machines roam to the strongest available signal and adds redundancy during carrier outages, while standardizing connectivity across Intelligent Machine Control, Smart Construction 3DMG, and Smart Quarry Site subscription services in the South Pacific. Autonomous mining equipment is excluded for now.

<https://www.itnews.com.au/news/komatsu-australia-reduces-reliance-on-telstra-locked-sims-in-security-push-628163>

THREAT INTELLIGENCE Check Point Software 2026-08-17

Check Point Research Weekly Report Flags Lazarus Zero-Day and Ransomware Surge

Check Point Research's weekly threat intelligence report detailed Lazarus-linked Operation Dream Job campaigns against defense organizations in Europe, India, and Brazil using fraudulent job postings, trojanized PDF software, and Windows zero-day CVE-2026-68820 to gain SYSTEM privileges and disable security visibility. The report counted 2,139 publicly reported ransomware victims in Q2 2026, up 33% year over year across 93 active groups, and averaged 2,336 weekly attacks per organization in July, with attackers increasingly using AI coding assistants.

<https://research.checkpoint.com/2026/17th-august-threat-intelligence-report/>

THREAT INTELLIGENCE

HiddenLayer 2026-08-17

HiddenLayer Research Shows Hidden Prompt Injections Can Hijack AI Coding Assistants Like Cursor

HiddenLayer published research detailing an indirect prompt injection chain against Cursor, in which hidden instructions planted in a benign-looking GitHub README coerce the AI coding assistant into exfiltrating API keys and SSH credentials and running blocked commands while in Auto-Run mode. The work also documents a denylist bypass, disclosed to Cursor and patched in version 1.3. HiddenLayer demonstrates its AI Detection and Response module blocking both direct and indirect injections in real time, positioning runtime AI protection as necessary for AI-assisted development.

<https://www.hiddenlayer.com/research/how-hidden-prompt-injections-can-hijack-ai-code-assistants-like-cursor>

FUNDING

Mindgard 2026-08-17

Mindgard's £22M Series A Fuels US Expansion for Automated AI Red Teaming

Coverage of Mindgard's £22M (roughly \$30M) Series A, led by Album VC with Karma Ventures and existing investors .406 Ventures, Atlantic Bridge, IQ Capital and Lakestar, bringing total funding to about \$42M. The Lancaster University spinout, headquartered in London and Boston, will scale product, engineering, sales and marketing while expanding in the US. Mindgard claims more than 150 publicly disclosed high-impact AI vulnerabilities across tools including ChatGPT, Google products and Cursor, and sells continuous automated red teaming plus runtime protection.

<https://businesscloud.co.uk/news/lancaster-uni-spinout-mindgard-raises-another-22m/>

COMPANY NEWS

Fastly 2026-08-17

Fastly Expands Revolving Credit Facility to \$100 Million and Extends Maturity to 2029

Fastly disclosed in an SEC filing that it executed a Fourth Amendment to its 2021 credit agreement, raising its senior secured revolving facility from \$60 million to \$100 million and extending maturity to August 17, 2029, with a possible one-year extension. Borrowing costs dropped 25 basis points to SOFR plus 1.75% or base rate plus 0.75%, with revised commitment fees. The amendment adds balance-sheet capacity as Fastly funds edge and security expansion.

<https://www.sec.gov/Archives/edgar/data/1517413/000151741326000238/fslly-20260817.htm>

THREAT INTELLIGENCE

Citrix (NetScaler) 2026-08-17

NetScaler ADC and Gateway Flaw CVE-2026-8452 Confirmed Exploited in the Wild

Threat intelligence tracking confirmed active exploitation of CVE-2026-8452, a memory overflow in NetScaler ADC and NetScaler Gateway, first asserted as exploited on August 17, 2026. The flaw stems from SAML signature canonicalization and can enable unauthenticated code execution or denial of service; a public proof-of-concept had already been released. Affected builds are 14.1 before 14.1-72.61 and 13.1 before 13.1-63.18, with no workaround available other than patching.

<https://db.gcve.eu/known-exploited-vulnerabilities-catalog/85bc654a-ccea-45ad-8fe9-ccd8adb5a39d>

ANALYST REPORT

Arista Networks 2026-08-17

Analysts Reset Arista Outlook on At Least \$3.6B in 2026 AI Networking Revenue

Analyst commentary highlighted Arista's upbeat AI networking outlook, with management guiding to at least \$3.6 billion of AI-related revenue in 2026 following Q2 revenue of \$3.04 billion and Q3 guidance near \$3.3 billion. The note flagged Ethernet-based AI fabric demand and upcoming OCP APAC Summit appearances in Taipei as near-term catalysts. Risks cited include heavy hyperscaler customer concentration and pricing pressure in high-performance switching.

<https://simplywall.st/stocks/us/tech/nyse-anet/arista-networks/news/the-bull-case-for-arista-networks-anet-could-change-followin-2>

PRODUCT LAUNCH

HPE / Juniper Networks 2026-08-17

HPE Positions PTX 12200 Router for Scale-Across AI Data Center Interconnect

HPE Networking detailed its PTX 12200 router as a high-scale data center interconnect platform for AI, built on Express 5 silicon with industry-leading 800G port density. Integrated ZR coherent optics simplify the optical layer and dark-fiber use, while line-rate MACsec provides encryption in transit. HPE framed the platform around scale-up, scale-out, and scale-across designs for distributed GPU clusters, with a roadmap toward 1.6T and 3.2T interfaces.

<https://convergedigest.com/hpes-ptx-12200-a-really-big-router-for-scale-across/>

COMPANY NEWS

Barracuda Networks 2026-08-17

Barracuda Integrated Email Protection Named Best Email Security in 2026 Tech Ascension Awards

Barracuda announced that Barracuda Integrated Email Protection was named Best Email Security Solution in the 2026 Tech Ascension Awards, citing AI-powered detection and rapid API-based deployment for Microsoft 365 and Google Workspace. The product combines behavioral AI, natural language processing, social graph analysis, and real-time URL inspection to catch phishing, business email compromise, account takeover, impersonation, and AI-assisted attacks. Barracuda pointed to customer results including removal of more than 12,000 malicious emails at one deployment.

<https://www.financialcontent.com/article/bizwire-2026-8-17-barracuda-named-best-email-security-solution-in-the-2026-tech-ascension-awards>

ACQUISITION

Fortinet 2026-08-17

Fortinet Acquires Virtue AI to Advance Continuous AI Protection

Fortinet announced the acquisition of Virtue AI, adding automated AI validation, runtime protection, and security for autonomous agentic systems to its AI-native Security Fabric. Virtue AI contributes Guardian Agent technology, automated red-teaming across hundreds of attack vectors and 1,000+ risk categories, and real-time guardrails that extend FortiAI Gate from development through runtime. Terms were undisclosed and described as immaterial, with enforcement and threat intelligence coordinated through FortiGuard Labs.

<https://www.fortinet.com/corporate/about-us/newsroom/press-releases/2026/fortinet-advances-continuous-ai-protection-with-the-acquisition-of-virtue-ai>

RESEARCH

CrowdStrike 2026-08-17

CrowdStrike Teaches Charlotte AI to Reason Through Detection Triage

CrowdStrike published research on a reasoning-enabled triage model that extends its NVIDIA Nemotron-powered Charlotte AI beyond fast true/false-positive verdicts to a transparent, step-by-step chain of reasoning analysts can audit. The classifier reasons across command lines and behavioral signals, reaching 82.6% accuracy and outperforming off-the-shelf and larger models. Initial coverage focuses on Windows endpoint detections, with plans to expand platforms, and the work supports CrowdStrike's Open Secure AI Alliance participation.

<https://www.crowdstrike.com/en-us/blog/teaching-ai-to-reason-through-detection-triage/>

PRODUCT LAUNCH

AWS 2026-08-17

AWS Ships EC2 Application Status Checks, IAM Role Manager, and OpenAI Daybreak on Bedrock

AWS's weekly roundup detailed new EC2 application-level status checks that detect app-layer failures such as a web server refusing requests or a stopped Docker daemon, plus an IAM role manager that auto-creates or reuses default roles across six service consoles. AWS also made OpenAI Daybreak models available to eligible customers on Amazon Bedrock, offering governed access to frontier AI for defensive security work and vulnerability research, alongside new foundation models in SageMaker JumpStart.

<https://aws.amazon.com/blogs/aws/aws-weekly-roundup-ec2-application-status-checks-iam-role-manager-openai-daybreak-on-bedrock-and-more-august-17-2026/>

COMPANY NEWS

Cisco 2026-08-17

Cisco's AI Networking Supercycle Puts It in Position to Sell Supercomputer Fabrics

Analysis of Cisco's record fourth quarter argued the company is finally positioned to sell large-scale AI supercomputer systems and the networks connecting them. Cisco posted \$17.3 billion in quarterly revenue with networking product orders up 40% year over year and \$4 billion in Q4 hyperscaler AI infrastructure orders, \$9.3 billion for fiscal 2026. Management guided to roughly \$7.5 billion of AI infrastructure revenue in fiscal 2027, signaling sustained hyperscaler and enterprise AI fabric demand.

<https://www.nextplatform.com/connect/2026/08/17/cisco-can-finally-sell-lots-of-supercomputers-and-their-networks/5288804>

COMPANY NEWS

Zscaler 2026-08-17

Zscaler CISO Calls for an Agentic SOC and a SecOps Revolution

Zscaler CISO Sam Curry argued that AI-driven threats require rebuilding security operations rather than automating existing workflows, urging teams to create new signals and instrumentation instead of merely correlating alerts. He described "agentic SOC" as a human-AI partnership that changes both outcomes and analyst development, favoring apprenticeship over rigid tiering. Curry also flagged destabilizing AI patches from frontier models and a broader, business-focused CISO mandate as emerging operational risks.

<https://www.databreachtoday.co.uk/ai-demands-secops-revolution-a-32565>

COMPANY NEWS Cloudflare 2026-08-17

Cloudflare Says Bot Blocking Is Driving New Publisher AI Licensing Deals

Cloudflare reported that its default AI crawler blocking and bot management controls are pushing AI companies to strike licensing agreements with publishers, turning unpaid scraping into negotiated access. The company positioned its bot detection and pay-per-crawl controls as leverage that restores publisher bargaining power over content used for model training. The framing extends bot management from a pure security control to a content monetization and access governance play.

<https://pressgazette.co.uk/platforms/cloudflare-says-bot-blocking-is-fueling-publisher-ai-deals/>

COMPANY NEWS Akamai 2026-08-17

Akamai Details Edge Architecture for Serving Humans and AI Crawlers Separately

Akamai published guidance on treating AI agents and human visitors as two distinct audiences, classifying each request at the edge with Bot Manager and then assembling either an AI-optimized or human-ready page using Akamai Functions, edge KV data, and GPUs. The approach caches both variants at the edge to cut origin trips, improve Core Web Vitals, and lower AI inference cost. It positions edge compute as the control point for agentic traffic rather than origin infrastructure.

<https://www.akamai.com/blog/trends/website-two-audiences-humans-ai>

COMPANY NEWS Akamai 2026-08-17

Akamai Makes the Case for Distributed Inference to Cut AI Agent Latency

Akamai SVP of Product Jon Alexander argued that AI workloads are moving from centralized training clusters toward multi-stage agent workflows that repeatedly call multiple models and tools across geographies. Each hop adds latency, so Akamai positions distributed inference closer to users and data as the architectural answer for real-time, multi-turn agents. The message frames AI traffic management and model-serving placement as a delivery problem, not just a compute problem.

<https://tfir.io/ai-agent-latency-distributed-inference-akamai/>

RESEARCH Akamai 2026-08-17

Akamai Advisory CISO Warns Agentic AI Shoppers Break Traditional Bot Detection

Akamai Advisory CISO Steve Winterfeld outlined how agentic AI shopping bots are shifting commerce traffic from human-driven to machine-to-machine, defeating bot-detection methods that rely on human behavioral cues. Akamai research cited commerce as the top target for API-related attacks, with 85% of surveyed commerce companies reporting API security incidents. The guidance points retailers toward API-layer controls and agent-aware policy rather than blanket bot blocking.

<https://tfir.io/agentic-ai-shoppers-api-security-commerce-akamai/>

COMPANY NEWS Fastly 2026-08-17

Fastly Scales India to 62 PoPs as AI Traffic Drives Edge Demand

Fastly detailed an aggressive India buildout, growing from 4 to 62 points of presence in two years and targeting a double-digit share of global business from what is now its fourth-largest market by potential. The company tied the expansion to AI traffic, which it said rose roughly 30% between January and May 2026, noting over 51% of AI requests reach origin infrastructure versus under 9% for human traffic. Fastly plans roughly 5x India go-to-market growth by end of 2027.

<https://brandequity.economictimes.indiatimes.com/news/digital/fastly-eyes-double-digit-india-contribution-as-ai-traffic-grows/133285821>

Barracuda Research Shows Microsoft Copilot Turning One Stolen Login Into \$247,500 Fraud

Barracuda researchers published a proof-of-concept showing how a single compromised Microsoft 365 session, with no malware, can be escalated through Microsoft Copilot into a \$247,500 wire fraud. Attackers used legitimate Copilot actions to suppress sign-in alerts, map the org and identify the CEO, locate a pending transfer, draft a convincing payment-change request, and set forwarding rules to intercept replies. Barracuda recommends phishing-resistant MFA and monitoring Copilot queries into invoices and executive mail.

<https://startupfortune.com/researchers-show-how-microsoft-copilot-can-turn-one-login-into-247500/>