

F5 Daily Briefing - Friday, August 28 - Sunday, August 30, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

27

ANNOUNCEMENTS

13

ACTIVE


5

ACTIVE CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (27)

PRODUCT LAUNCH **Cloudflare** Sun, Aug 30, 2026 AI Data Delivery

Cloudflare AI Search Adds GLM-5.3 Flash With 1M-Token Context Window

Cloudflare's changelog announced GLM-5.3 Flash support in AI Search, running on Workers AI with a 1,048,576-token context window. Customers can select the model for their AI Search instance through the supported-models configuration. The addition continues Cloudflare's pattern of rapidly onboarding new inference models to its edge AI serving stack, reinforcing its pitch that model choice and long-context retrieval should live on the delivery network rather than in a separate AI platform.

<https://developers.cloudflare.com/changelog/post/2026-08-30-glm-5.3-flash/>

PRODUCT LAUNCH **AWS** Sun, Aug 30, 2026 Platforms

AWS Open Sources Kiro Crew for Asynchronous Coding Agents

AWS released Kiro Crew as open source, a framework for running asynchronous coding agents. The move extends AWS's agentic developer tooling into the open ecosystem, letting teams orchestrate multiple long-running AI coding agents rather than single synchronous sessions. Open sourcing signals AWS is competing for developer mindshare in agent orchestration standards, which increasingly determines where agentic workloads and their supporting inference traffic land.

<https://www.infoq.com/news/2026/08/kiro-crew-coding-agents/>

PARTNERSHIP **Fortinet** Sat, Aug 29, 2026 Platforms

Fortinet Signs Multi-Year Stanford Athletics Sponsorship With Stadium Naming Rights

Fortinet became an official sponsor of Stanford Athletics under a multi-year agreement starting in the 2026-27 academic year. The deal includes naming entitlement for the Fortinet Director's Level at Stanford Stadium, endzone and on-court branding at Maples Pavilion, and in-venue signage across football, basketball, baseball, soccer, softball, and volleyball. Fortinet framed the sponsorship as reinforcing its Silicon Valley roots and Bay Area brand presence, extending its consumer-visible marketing footprint well beyond traditional enterprise security channels.

<https://www.manilatimes.net/2026/08/29/tmt-newswire/globenewswire/stanford-athletics-and-cybersecurity-leader-fortinet-team-up-for-multi-year-sponsorship/2414325>

RESEARCH/REPORT **CrowdStrike** Sat, Aug 29, 2026 Platforms

theCUBE Research Analyzes CrowdStrike's Post-Mythos Surge and Blast-Radius Risk

theCUBE Research published a Breaking Analysis on CrowdStrike's momentum following its Mythos launch, examining the durability of its competitive moat and what analysts call the blast-radius test. The piece credits strong Falcon Flex adoption and accelerating ARR growth while positioning CrowdStrike's AI-agent containment approach, spanning identity control plane, data protection, runtime visibility, and guardrails, as the basis of its next expansion phase. It flags the upcoming Fal.Con event as the checkpoint for validating pipeline durability and continued product innovation.

<https://thecubereseearch.com/325-breaking-analysis-crowdstrikes-post-mythos-surge-moat-momentum-and-the-blast-radius-test/>

RESEARCH/REPORT **Cloudflare** Sat, Aug 29, 2026 Platforms

Cloudflare for Government Earns FedRAMP High, Opening More Federal Workloads

Cloudflare for Government received FedRAMP High authorization, clearing federal agencies to run highly sensitive national security, critical infrastructure, and financial workloads on the platform. More than 100 federal agencies already use Cloudflare, including State, Justice, Homeland Security, Energy, and Commerce. Coverage cited a \$7.7M five-year Magic Transit and Network Firewall contract with a large agency, and noted Q2 revenue of \$696.1M, up 35.9%, with full-year guidance raised.

<https://247wallst.com/investing/2026/08/29/the-next-palantir-100-federal-agencies-already-use-cloudflare-and-its-cleared-for-more/>

PRODUCT LAUNCH **Cloudflare** Sat, Aug 29, 2026 Application Delivery (ADCs)

Cloudflare Workers Now Accept Inbound TCP, Starting With gRPC

Cloudflare Workers can now terminate raw inbound TCP connections, with gRPC as the first supported protocol on top. The change lets developers front gRPC services directly on Cloudflare's edge rather than routing through HTTP-only entry points, extending Workers from an HTTP proxy layer toward a general-purpose ingress and reverse-proxy tier. It positions Workers to handle service-to-service and microservice traffic patterns that previously required dedicated load balancers or ingress controllers.

<https://www.infoq.com/news/2026/08/workers-inbound-tcp-grpc/>

RESEARCH/REPORT **AWS** Sat, Aug 29, 2026 Platforms

AWS Networking Efficiency and Interconnect Pricing Shift Draw Analyst Attention

The Register argued AWS is underselling its networking economics: a flat, single-tier, near-random fabric design plus its Resilient Network Graph approach delivers up to 40% better network energy efficiency and is now default for new data center builds. AWS also introduced AWS Interconnect for multicloud traffic with hourly port charges instead of per-GB fees, plus a 500 Mbps free tier per provider for connections to Google Cloud, Oracle, and soon Azure.

https://www.theregister.com/networks/2026/08/29/aws-mumbles-about-its-cost-busting-networking-tech-when-it_should_be_shouting/5293379

PRODUCT LAUNCH **AWS** Sat, Aug 29, 2026 AI Data Delivery

NVIDIA Cosmos 3 World Models Now Available on Amazon SageMaker

AWS said NVIDIA's Cosmos 3 world models are now available on Amazon SageMaker, adding physical-AI and simulation foundation models to its managed ML platform. The listing gives robotics, autonomous systems, and industrial simulation teams a managed path to fine-tune and serve world models on AWS infrastructure, deepening the AWS and NVIDIA model distribution relationship beyond raw GPU capacity into curated model serving.

<https://pivotnews.ai/five/aws-says-nvidia-s-cosmos-3-world-models-are-now-on>

THREAT INTEL **Zscaler** Fri, Aug 28, 2026 Platforms

Zscaler Report Finds Mobile, IoT, and OT Threat Activity Surging

Zscaler released research showing sharp growth in threats across mobile, IoT, and operational technology environments. Financially motivated mobile threats rose 111% year over year even as overall Android attacks declined, while IoT interactions grew 37% and IoT malware climbed 45%, with routers absorbing more than 65% of IoT attacks and Mirai and Gafgyt dominating malware families. Roughly half of observed OT devices still run end-of-life operating systems with exposed protocols, which Zscaler uses to argue for zero-trust segmentation across unmanaged and legacy infrastructure.

<https://www.swapupdate.in/zscaler-report-mobile-iot-and-ot-cyber-threats-surge/>

RESEARCH/REPORT **Palo Alto Networks** Fri, Aug 28, 2026 AI Security

Unit 42 Introduces Perturbation Probing to Expose Fragility in LLM Safety Alignment

Palo Alto Networks' Unit 42 published research introducing Perturbation Probing, a two-pass-per-prompt diagnostic that inspects feed-forward network behavioral circuits inside aligned large language models to reveal how easily safety guardrails break. Paired with the team's Logit-Gap Steering work, the method shows model-level alignment alone is unreliable and argues for defense-in-depth at runtime. Palo Alto positions Prisma AIRS Runtime Security as the enforcement layer, applying external content filtering and inline guardrails around deployed models.

<https://unit42.paloaltonetworks.com/perturbation-probing-llm-safety/>

RESEARCH/REPORT **Palo Alto Networks** Fri, Aug 28, 2026 AI Security

Palo Alto Unit 42 Warns Frontier AI Is Tipping the Scales Toward Cyber Adversaries

Unit 42 researchers warned that frontier AI models are accelerating attacker capability faster than defenses can adapt, letting a single operator automate vulnerability discovery and exploitation at near nation-state scale. In one exercise, AI identified 50 vulnerabilities and attack paths in 10 hours, work that would take humans roughly two weeks. Palo Alto estimates a shrinking three-to-five-month window before attackers routinely operate at that speed, and cited its participation in Anthropic and OpenAI AI security initiatives as part of a collective defense response.

<https://www.cybersecuritydive.com/news/frontier-ai-tipping-scales-cyber-adversaries/829088/>

PRODUCT LAUNCH **Operant AI** Fri, Aug 28, 2026 AI Security

Operant AI Launches Semantic Firewall to Enforce AI Agent Intent in Real Time

Operant AI released its Semantic Firewall, a runtime control plane that reads agent intent across prompts, tool calls, and data movement and makes inline allow, block, or redact decisions. It ships with Tool Intent Guard, Code Intent Guard, Data Intent Guard, and Scope Guard, plus Token Meter for per-session cost caps, Browser AI Protection, and expanded Claude coverage. Trade coverage frames it as a response to enterprises scaling autonomous agents faster than their governance controls, with emphasis on enterprise-owned audit telemetry rather than third-party model safeguards.

<https://cxotoday.com/media-coverage/operant-ai-launches-semantic-firewall-to-enforce-ai-agents-in-real-time-as-enterprise-adoption-scales-in-india/>

PRODUCT LAUNCH **HPE / Juniper Networks** Fri, Aug 28, 2026 AI Data Delivery

HPE Showcases Sovereign AI Factory With NVIDIA and Autonomous Networking at LEAP 2026

HPE brought its AI Factory with NVIDIA portfolio, autonomous networking from HPE Aruba and Juniper, and the Cray GX5000 with direct liquid cooling to LEAP 2026 in Riyadh. The sovereign AI positioning targets government and regulated customers needing in-country model training and inference, paired with self-driving network operations. HPE also featured Private Cloud AI, cyber resilience, and AI-ready data and storage, aligning the pitch with Saudi Vision 2030 digital sovereignty goals.

<https://intlbm.com/2026/08/28/hpe-brings-core-innovations-including-sovereign-ai-factory-with-nvidia/>

PRODUCT LAUNCH **HPE / Juniper Networks** Fri, Aug 28, 2026 AI Data Delivery

HPE Expands 'Saudi Made' Server Portfolio With ProLiant Gen12 Launch

HPE expanded its locally manufactured Saudi Made lineup with ProLiant Compute DL360 Gen12 and DL380 Gen12 servers, built at alfanar's Riyadh facility on Intel Xeon 6 processors. HPE cites up to 41% better performance per watt and up to 65% annual power savings, targeting hybrid cloud, analytics, virtualization, and AI workloads. Availability spans Saudi Arabia plus Jordan, Egypt, and GCC markets, reinforcing the in-country production and data sovereignty angle.

<https://channelpostmea.com/2026/08/28/hpe-expands-its-saudi-made-server-portfolio-with-new-next-generation-server-launch/>

RESEARCH/REPORT **Fastly** Fri, Aug 28, 2026 AI Data Delivery

Fastly's Simon Wistow on the Bot-First Web, Poisoned Crawlers, and Why Edge GPUs Don't Add Up

In an iTWire feature, Fastly co-founder Simon Wistow argued the web is becoming machine-first, with AI agent and crawler traffic growing far faster than human traffic and increasingly bypassing cache to hit origin. He described poisoned-crawler tactics and pushed back on the idea of putting GPUs at the edge, saying inference economics favor centralized accelerators with the edge handling identification, routing, and control of machine traffic instead.

<https://itwire.com/it-industry-news/telecoms-and-nbn/feature-fastlys-simon-wistow-on-the-bot-first-web-poisoned-crawlers-and-why-gpus-at-the-edge-dont-add-up>

PARTNERSHIP **CrowdStrike** Fri, Aug 28, 2026 Platforms

CrowdStrike and Telkom Indonesia Sign MoU to Advance AI-Powered Cybersecurity

CrowdStrike signed a Memorandum of Understanding with PT Telkom Indonesia to deliver AI-driven protection and managed security services across the Indonesian market. Telkom plans to standardize its own cyber defenses on the Falcon platform, integrate Falcon with local capabilities, and build joint go-to-market offerings. The two companies will also explore co-developing solutions to secure AI agents, models, and systems for enterprise AI adoption, alongside managed security services and quantum-resilience initiatives, extending CrowdStrike's telco-led sovereign and service-provider distribution motion in Southeast Asia.

<https://www.crowdstrike.com/en-us/press-releases/crowdstrike-and-telkom-indonesia-enter-into-mou-to-advance-ai-powered-cybersecurity/>

EARNINGS **CrowdStrike** Fri, Aug 28, 2026 Platforms

CrowdStrike Q2 Earnings Call Centers on the Expanding Enterprise AI Attack Surface

On its earnings call CrowdStrike framed a structural shift in security driven by autonomous AI models and non-human digital identities. CEO George Kurtz described an arms race in which AI accelerates both attacks and security spending, pushing customers to consolidate onto platforms and extend protection past traditional endpoints. Management pointed to rising demand for integrated defenses spanning endpoint, cloud, and machine identities as autonomous agents enter production workflows, arguing that machine-speed, continuous protection is becoming a baseline enterprise requirement rather than a premium add-on.

<https://securitytoday.com/articles/2026/08/28/crowdstrike-earnings-highlight-expanding-enterprise-ai-attack-surface.aspx>

PRODUCT LAUNCH **Cloudflare** Fri, Aug 28, 2026 Platforms

Cloudflare One Client Reaches GA on Windows, macOS, and Linux

Cloudflare shipped generally available Cloudflare One Client releases for Windows (2026.7.1376.0), macOS (2026.7.1376.0), and Linux (2026.7.1377.0) on the same day. The coordinated GA across all three desktop platforms keeps the WARP-based access client aligned for enterprises standardizing on Cloudflare One for zero trust connectivity, reducing version drift for IT teams managing mixed endpoint fleets.

<https://developers.cloudflare.com/changelog/post/2026-08-28-warp-windows-ga/>

RESEARCH/REPORT **Cloudflare** Fri, Aug 28, 2026 Platforms

Cloudflare Frees 100TB of RAM by Shrinking 1.1.1.1 DNS Cache Entries

Cloudflare detailed an engineering effort that trimmed its 1.1.1.1 resolver cache entries by tens of bytes each, reclaiming roughly 100TB of memory across its fleet. With about 250 billion cached DNS entries at any moment, a single wasted byte costs 250GB, making micro-optimizations material at Cloudflare's scale. The work is a proof point for the efficiency-of-scale argument Cloudflare uses when positioning its global network against smaller delivery and DNS providers.

<https://www.tomshardware.com/tech-industry/big-tech/cloudflare-frees-100tb-of-ram-by-shrinking-dns-cache-entries>

PRODUCT LAUNCH **Cloudflare** Fri, Aug 28, 2026 Application Security

Cloudflare Launches BotBase for Operators to Register and Track Bots and AI Agents

Cloudflare expanded BotBase from a bot directory for website owners into a self-service channel for bot operators. Operators can now submit their bots, describe behavior, content use, and identity, then view submission status and history in a new BotBase section under Protect & Connect, Application Security. Cloudflare frames it as a transparency step toward verified identity, observability, and more polite crawling, with automated review now handling a submission volume roughly 7x higher than 2023.

<https://blog.cloudflare.com/botbase-for-operators/>

THREAT INTEL **Citrix (NetScaler)** Fri, Aug 28, 2026 Application Delivery (ADCs)

NetScaler Flaw CVE-2026-8452 Hits CISA KEV as Federal Patch Deadline Lands

CISA's Known Exploited Vulnerabilities update put Citrix NetScaler ADC and NetScaler Gateway back in the spotlight, with CVE-2026-8452, an improper memory-buffer restriction enabling unauthenticated remote code execution, among six newly added exploited flaws. Federal agencies faced an August 29 remediation deadline under BOD 26-04, while the other five legacy flaws carry a September 9 date. Reports cite in-the-wild exploitation against Gateway/AAA vserver and SSL VPN configurations, keeping NetScaler estates under urgent upgrade pressure.

<https://www.secure.com/news/cisa-kev-six-flaws-netscaler-linux-sql-server>

PRODUCT LAUNCH **Cisco** Fri, Aug 28, 2026 Platforms

Cisco Rolls Out MyAgent AI Agent for Autonomous Enterprise Workflows to 90,000 Employees

Cisco began deploying MyAgent, an AI agent that moves beyond chat into supervised autonomous execution across Outlook, Webex, Jira, and SharePoint, starting with its own 90,000 employees. Built on Cisco's Circuit platform for governed multi-model AI, MyAgent restricts activity to approved LLMs and enterprise data paths, accepts intent-based goals, plans its own steps, and retains persistent memory for personalization. Cisco describes it as an ambient execution layer with human oversight built in, an internal proof point for its agentic enterprise platform story.

<https://www.electronicmedia.info/2026/08/28/cisco-myagent/>

RESEARCH/REPORT Cisco Fri, Aug 28, 2026 AI Security

Cisco Research Says AI Model Country Labels Are Meaningless Without Lineage Tracing

Cisco researchers, working with VAIL, reported that country-of-origin labels on AI models fail to reflect what is actually inside them, because weights and training dependencies are inherited across publishers. Using a Model Provenance Kit and behavioral fingerprinting, the team found substantial lineage ties between models even after rebranding. Cisco proposes a model bill of materials documenting base checkpoints, derivation methods, datasets, synthetic data tools, and licenses, warning enterprises that banning a vendor by nationality does not remove the underlying model dependencies or embedded biases.

<https://www.securityweek.com/think-youve-eliminated-chinese-ai-check-the-models-lineage-cisco-says/>

THREAT INTEL Barracuda Networks Fri, Aug 28, 2026 Application Security

Barracuda Research Details Callback Phishing Using Fake Invoices and Calendar Invites

Barracuda researchers published an analysis of callback phishing, or TOAD (telephone-oriented attack delivery), showing how attackers use fake invoices, billing receipts and calendar invites to push victims into calling attacker-controlled support lines. The report walks through five campaigns impersonating brands such as PayPal and Geek Squad, with payload text hidden in images and attachments to evade conventional email filters and ICS invites used to bypass inbox controls. Barracuda argues defenses need OCR-based image analysis, impersonation and language checks, and coverage across both email and collaboration channels.

<https://blog.barracuda.com/2026/08/28/callback-phishing-fake-invoice-toad-attacks>

PRODUCT LAUNCH Akamai Fri, Aug 28, 2026 Application Delivery (ADCs)

Akamai Property Manager Moves gRPC Support to General Availability

Akamai announced general availability of gRPC in Property Manager, including new gRPC-Web text translation support. Customers need to confirm gRPC is included in their contract with their account representative. The GA lets Akamai front modern gRPC-based microservice and API traffic on its delivery platform rather than requiring HTTP/1-style workarounds, closing a protocol gap against edge platforms already handling gRPC ingress.

<https://techdocs.akamai.com/property-mgr/changelog/aug-28-2026-grpc-in-general-availability>

THREAT INTEL Akamai Fri, Aug 28, 2026 Application Security

Akamai Details Zero-Day Defense Ahead of Public Vulnerability Disclosure

Akamai published a security post describing how it protects customers against zero-day vulnerabilities before the industry discloses them. The piece outlines rapid threat detection, staged protection rollouts, asset-specific mitigations, and shared intelligence across its platform and threat research team. Akamai positions preemptive, platform-delivered virtual patching as an advantage over customer-side patch cycles, arguing early protection materially reduces exposure windows during the gap before CVEs become public.

<https://www.akamai.com/blog/security/zero-day-akamai-defends-before-industry-discloses>

PRODUCT LAUNCH AWS Fri, Aug 28, 2026 AI Data Delivery

Amazon EC2 P6-B300 GPU Instances Expand to Additional AWS Regions

AWS made its P6-B300 accelerated compute instances available in more regions, widening the geographic footprint for large-scale AI training and inference workloads. The expansion gives customers closer placement options for GPU capacity, reducing latency and helping address data residency requirements for AI serving. It reflects AWS's continued push to make high-end accelerator capacity a regionally distributed utility rather than a scarce resource concentrated in a few hub regions.

<https://aws.amazon.com/about-aws/whats-new/2026/08/amazon-ec2-p6-b300-instances-available-additional-regions/>