

F5 Daily Briefing - Thursday, September 03, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

29


ANNOUNCEMENTS

15

ACTIVE

5

ACTIVE CATEGORIES



ALL ANNOUNCEMENTS (29)

ACQUISITION

Palo Alto Networks

Sep 03, 2026

Palo Alto Networks Acquires AI Agent Startup Console for \$500M

Palo Alto Networks confirmed a roughly \$500 million cash-and-stock acquisition of Console, a 2024-founded startup whose AI agents automate password resets, permission provisioning, and IT troubleshooting for customers including Ramp, Flock Safety, and Scale AI. Console folds into the Cortex platform, adding natural-language, agent-driven investigation and remediation of security alerts. This is Palo Alto's seventh acquisition of 2026, following Chronosphere and Koi, and extends its platformization push deeper into autonomous security operations and adjacent IT workflows.

<https://pulse2.com/palo-alto-networks-acquires-console-to-expand-agentic-security-capabilities/>

EARNINGS

Palo Alto Networks

Sep 03, 2026

Palo Alto Networks Earnings Highlight SASE and Network Security Wins

Coverage of Palo Alto Networks' latest results spotlights continued momentum in SASE and network security as the primary growth engines behind its platformization strategy. The report frames large multi-product platform deals and consolidation displacement as the key drivers, with SASE attach rates cited as evidence customers are collapsing point tools onto a single vendor. The results give a current read on how effectively platform bundling is converting competitive displacement into recurring revenue.

<https://www.sdxcentral.com/news/palo-alto-networks-earnings-highlight-sase-and-network-security-wins/>

BLOG POST**Palo Alto Networks**

Sep 03, 2026

Orchestrating the Agentic SOC: Palo Alto Networks Details Model Context Protocol Strategy

Palo Alto Networks published a practical guide positioning Model Context Protocol as a standardized AI integration layer connecting large language models to SOC tooling. The post details the Cortex MCP Server and native MCP integrations spanning Cortex XSIAM, AgentiX, XDR, and Cloud, letting agents orchestrate tools on demand in response to analyst prompts or events. The stated goal is eliminating the 'integration tax' of brittle custom scripts and enabling agent-driven, end-to-end AI security operations.

<https://www.paloaltonetworks.com/blog/security-operations/mcp-client/>

PARTNERSHIP**Fortinet**

Sep 03, 2026

Fortinet and FIRST: Strengthening Cyber Resilience Through Global Collaboration

Fortinet detailed its expanded partnership with FIRST, the global Forum of Incident Response and Security Teams, which it has worked with since 2012 through FortiGuard Labs threat intelligence and national CERT integrations. As a founding partner of FIRST CORE, Fortinet is funding capacity building, training, and regional exercises in under-resourced regions. Stated priorities include expanding engagement in Africa, adding training programs, hosting more regional symposia, and collaborating with FIRST Special Interest Groups on AI security and ransomware response.

<https://www.fortinet.com/blog/industry-trends/fortinet-and-first-strengthening-cyber-resilience-through-global-collaboration>

SECURITY ADVISORY**Fortinet**

Sep 03, 2026

Fortinet Links FortiBleed Credentials to Prior Incidents

New reporting says Fortinet has traced credentials used in the ongoing FortiBleed campaign against FortiGate devices back to earlier disclosed incidents rather than a new vulnerability, with attackers brute-forcing appliances that have weak passwords and no MFA. Fortinet's guidance centers on resetting admin and VPN credentials, enforcing MFA, upgrading to releases using stronger password hashing, and restricting internet-exposed management interfaces. The campaign has been tied to ransomware activity, keeping edge-device credential hygiene in the spotlight.

<https://insight.tmcnet.com/insight/fortinet-traces-fortibleed-credentials-to-earlier-incidents-mtlgn7qe>

SECURITY ADVISORY**Cisco**

Sep 03, 2026

Cisco Warns of Unpatched Secure Email Flaws, Patches Critical Switch Vulnerabilities

Cisco disclosed a critical unauthenticated remote code execution flaw in Nexus 9000 Series switches running Silicon One (CVE-2026-20212, CVSS 9.8) alongside an IOS XR security hardening release, while warning that ciphertext-decryption weaknesses in Cisco Secure Email remain unpatched. Interim mitigations include infrastructure ACLs limiting management and control-plane access plus a Live Protect shield for select NX-OS releases. Cisco reported no known exploitation at disclosure but noted the window between disclosure and exploitation continues to shrink.

<https://www.securityweek.com/cisco-warns-of-unpatched-secure-email-flaws-patches-critical-switch-vulnerabilities/>

RESEARCH REPORT**Cisco**

Sep 03, 2026

Your Agent Trusts Things You Never Approved: Cisco Details AI Defense Agent Supply-Chain Scanning

Cisco published research on agentic AI supply-chain trust, arguing that agents load skills and MCP servers from lightly vetted registries at machine speed. It describes a partnership with OpenAI bringing Daybreak models into Cisco AI Defense as a tiered scanner: Daybreak Blue for a cybersecurity-tuned first pass and Daybreak Red for heavily obfuscated payloads and exploit chains. Cisco's skill-scanner and mcp-scanner combine YARA, static analysis, Python bytecode checks, shell taint analysis, and semantic review before trust is granted.

<https://blogs.cisco.com/ai/your-agent-trusts-things-you-never-approved>

PRODUCT LAUNCH**CrowdStrike**

Sep 03, 2026

CrowdStrike Unveils SafeMind for Agentic Cyber

Analyst coverage of CrowdStrike's Fal.Con announcements examines SafeMind, a family of purpose-built security models and harnesses built with NVIDIA that runs natively in the Falcon platform. SafeMind pairs offensive (Red Tempest) and defensive (Blue Solano) models in one closed-loop agentic system that continuously finds and closes attack paths, trained on Falcon telemetry. CrowdStrike claims materially faster end-to-end remediation and lower detection cost versus general-purpose frontier models, positioning owned full-stack models as a differentiator against generic AI.

<https://aragonresearch.com/crowdstrike-unveils-safemind-for-agentic-cyber/>**PRODUCT LAUNCH****CrowdStrike**

Sep 03, 2026

CrowdStrike Launches AI Security Tools for Enterprises

Reporting from Fal.Con 2026 summarizes CrowdStrike's enterprise AI security push, including agentic SOC capabilities, an Agentic Identity Provider for AI agent identity, Falcon Guardian AI detection and response, and supply-chain protections. The set targets runtime monitoring of autonomous agents across endpoints, SaaS, cloud, and browsers, plus governance of agent identities. It underscores how quickly platform vendors are packaging AI-agent protection as a standard tier rather than a standalone product.

<https://securitybrief.com.au/story/crowdstrike-launches-ai-security-tools-for-enterprises>**PARTNERSHIP****Zscaler**

Sep 03, 2026

Zscaler and OpenAI Bring Frontier AI-Powered Security to the Enterprise

Zscaler unveiled two new capabilities built with OpenAI's Daybreak Defense Network and GPT cyber models. Endpoint AI Security reasons over full endpoint state — including AI agents, coding assistants and IDE/browser extensions — to reconstruct realistic multi-stage attack chains instead of flat CVE lists. Identity Risk Analysis in the Zscaler AI Access Graph traces risky permission paths across humans, service accounts and AI agents spanning Azure, AWS, OCI and GCP, ranking exposure by business impact. Zscaler stresses evidence-grounded output with human review before consequential actions.

<https://www.zscaler.com/blogs/product-insights/zscaler-and-openai-bringing-frontier-ai-powered-security-enterprise>**EARNINGS****Zscaler**

Sep 03, 2026

Zscaler Announces Strong Fourth Quarter and Fiscal 2026 Financial Results

Zscaler reported fiscal Q4 and full-year 2026 results for the period ended July 31, 2026, with revenue and ARR each up 25% year over year and net-new ARR up 24%. Non-GAAP operating margin reached a record 24%. Management credited broadening growth drivers beyond seat-based user solutions — including non-seat-based products, Z-Flex momentum, larger deal activity and improved sales productivity — and guided to continued momentum in fiscal 2027, with AI agent security positioned as a next growth driver.

<https://ir.zscaler.com/news-releases/news-release-details/zscaler-announces-strong-fourth-quarter-and-fiscal-2026>**PARTNERSHIP****Check Point Software**

Sep 03, 2026

Check Point Brings OpenAI Daybreak Models Across Its Security Platform

Check Point is embedding OpenAI's Daybreak frontier cyber models across its security platform through the Daybreak Defense Network. Targeted workflows include Exposure Validation to prove exploitable risk and prioritize remediation, Keystone for intent-driven autonomous policy management, the Autonomous Workspace Platform for faster incident investigation, and NexPloit to convert vulnerability data into verified attack material that speeds protection development. The rollout is phased, with governance guardrails limiting model visibility, constraining actions and validating outputs before automation.

<https://blog.checkpoint.com/ai-security/check-point-brings-openai-daybreak-models-across-its-security-platform-to-help-defenders-find-validate-and-remediate-risk/>

PRODUCT LAUNCH**Cloudflare**

Sep 03, 2026

Cloudflare Partners with OpenAI Daybreak Models to Redefine Vulnerability Management with AI-Powered Edge Defense

Cloudflare launched Vulnerability Discovery and Remediation, an invitation-only early access service inside Cloudflare Managed Defense powered by OpenAI Daybreak models including GPT-5.6 Cyber. The service correlates live edge traffic with reconnaissance and hunting across customer-authorized codebases, then deploys WAF rules at the edge for immediate virtual patching and generates verified code patches for developer review. It spans Cloudflare Web Assets, WAF and Workers Observability. Inference runs through Cloudflare AI Gateway to OpenAI rather than on the edge, and every proposed fix or rule requires customer approval.

<https://www.cloudflare.com/press/press-releases/2026/cloudflare-partners-with-openai-daybreak-models-to-redefine-vulnerability-management-with-ai-powered-edge-defense/>

PRODUCT LAUNCH**AWS**

Sep 03, 2026

Amazon Linux 2027 Is Now Available in Public Preview

AWS opened public preview of Amazon Linux 2027, built on AL2023 for cloud-native workloads with a focus on performance, scale and security. It ships kernel 7.1+ with SELinux enforcing by default, accelerated cryptography via AWS-LC, refreshed toolchains and language runtimes, and AI/ML support including AWS Neuron driver access for accelerated inference and training instances. Preview AMIs are available across all commercial AWS Regions for x86-64 and ARM, with container base images on the Amazon ECR Public Gallery.

<https://aws.amazon.com/about-aws/whats-new/2026/09/announcing-amazon-linux-2027/>

PRODUCT LAUNCH**AWS**

Sep 03, 2026

Amazon S3 Now Supports PrivateLink for FIPS Endpoints

AWS extended AWS PrivateLink support to Amazon S3 FIPS endpoints, letting regulated customers reach S3 over private connectivity while meeting FIPS 140 cryptographic module requirements. Traffic stays on the AWS network instead of traversing the public internet, simplifying compliance architectures for government, defense and financial workloads that previously had to choose between private connectivity and FIPS-validated endpoints.

<https://aws.amazon.com/about-aws/whats-new/2026/09/amazon-s3-privatelink-fips-endpoints/>

PRODUCT LAUNCH**AWS**

Sep 03, 2026

AWS Gateway Load Balancer Now Supports TCP Reset for Faster Failure Recovery

AWS Gateway Load Balancer can now send TCP Reset packets to accelerate failure recovery for inline security appliance deployments. RSTs are issued when a target becomes unhealthy, when a target is deregistered after connection draining, and when the TCP idle timeout expires — cutting traffic interruption from minutes to seconds by signaling invalid connections to clients and servers immediately. The behavior is off by default for backward compatibility and is enabled per target group via Console, CLI or API, at no additional cost in all Regions where GWLB is available.

<https://aws.amazon.com/about-aws/whats-new/2026/09/aws-gateway-load-balancer-tcp-reset/>

PRODUCT LAUNCH**AWS**

Sep 03, 2026

Amazon CloudFront Announces API Support for Flat-Rate Pricing Plans

Amazon CloudFront added API support for its flat-rate pricing plans, letting customers programmatically select, provision and manage predictable-cost CDN tiers instead of configuring them only through the console. The change makes flat-rate delivery pricing consumable through infrastructure-as-code and automated account provisioning, reinforcing AWS's push to simplify CDN cost predictability for content and application delivery buyers.

<https://aws.amazon.com/about-aws/whats-new/2026/09/cloudfront-flat-rate-pricing-plans-api/>

PRODUCT LAUNCH**AWS**

Sep 03, 2026

Introducing Cross-Region Inference for OpenAI GPT-5.6 Models on Amazon Bedrock

Amazon Bedrock now offers OpenAI GPT-5.6 models across 25+ AWS Regions using cross-Region inference profiles (CRIS). Three variants — Sol, Terra and Luna — balance capability and cost, accept text and image inputs, and support a 1 million token context window with reasoning modes. CRIS automatically routes requests from a source Region to supported destination Regions, expanding usable capacity and absorbing traffic bursts without customers managing routing themselves. Models are reachable through OpenAI-compatible Responses, Chat Completions and Converse APIs.

<https://aws.amazon.com/blogs/machine-learning/introducing-cross-region-inference-for-openai-gpt-5-6-models-on-amazon-bedrock/>

FUNDING**HiddenLayer**

Sep 03, 2026

HiddenLayer Raises \$100 Million Series B for AI Runtime Security

HiddenLayer closed a \$100 million Series B led by Delta-v Capital with Ten Eleven Ventures, Morgan Stanley, Microsoft's M12, and Booz Allen Ventures participating. The Austin-based firm says ARR grew more than 10x in the past year and will use the capital to expand Agentic Runtime Security and a new Agent Harness Security offering that protects AI coding agents in production. The round funds international go-to-market expansion and deepens coverage of prompt injection, agent manipulation, and malicious tool use.

<https://www.securityweek.com/hiddenlayer-raises-100-million-for-ai-runtime-security/>

PRODUCT LAUNCH**Lasso Security**

Sep 03, 2026

Lasso Security Launches LEAP CPU-Based AI Guardrails and Raises \$30M

Lasso Security unveiled LEAP, a transformer-free AI guardrail engine that runs on ordinary CPUs with sub-5ms decision latency and no GPU footprint, claiming GPU-class detection accuracy at thousands of times the throughput of existing guardrails. LEAP is already in production at global enterprises and the U.S. federal government and pairs with RAPID, Lasso's self-hosted model for nuanced judgments, behind a routing layer. Lasso also disclosed a \$30 million round led by ClearSky with Entrée Capital, iAngels, Singtel Innov8, Mindset, and Swish Data.

<https://www.lasso.security/resources/lasso-security-announces-future-of-ai-security-with-cpu-based-guardrails>

FUNDING**Mindgard**

Sep 03, 2026

Mindgard Confirms \$30M Series A to Scale Automated AI Red Teaming

Lancaster University confirmed that spin-out Mindgard has raised a \$30 million (£22.2m) Series A led by Album VC, with Karma Ventures and existing backers .406 Ventures, Atlantic Bridge, IQ Capital, and Lakestar. It is among the largest Series A cybersecurity rounds for a UK company and will scale product engineering, sales, and marketing. Mindgard combines AI safety research with offensive security to red-team models, agents, and applications, with deployments across financial services, pharma, gaming, and healthcare.

<https://www.lancaster.ac.uk/news/major-funding-boost-for-lancaster-ai-spin-out-company>

PRODUCT LAUNCH**Akamai**

Sep 03, 2026

Akamai Introduces More Granular Controls for AI Bot Traffic

Akamai split its single "AI Bots" category in the Akamai Bot Directory into three distinct types — AI training crawlers, AI search crawlers, and AI fetchers/agents — so customers can write policy beyond blanket block/allow. Teams can now block training scrapers while still permitting search indexers, and distinguish real-time agentic user traffic from background crawlers. The controls are enforced through Bot & Agent Control and pair with Akamai's AI Brand Presence capability to preserve visibility in AI search results while protecting intellectual property.

<https://www.akamai.com/blog/security/introducing-more-granular-controls-ai-bot-traffic>

RESEARCH REPORT**Fastly**

Sep 03, 2026

Fastly Named a 2026 Gartner Peer Insights Customers' Choice for Edge Distribution Platforms

Fastly announced it was named a 2026 Gartner Peer Insights Customers' Choice for Edge Distribution Platforms, based on 92 verified reviews with a 4.8/5 overall rating and 95% willingness to recommend as of April 30, 2026. The recognition cites platform performance, reliability, customer experience, and support. It follows seven consecutive years of Customers' Choice recognition for Cloud Web Application and API Protection (WAAP), and Fastly framed the award around enterprises scaling AI workloads on its edge platform.

<http://www.upgrademag.com/web/2026/09/03/fastly-named-a-2026-gartner-peer-insights-customers-choice-for-edge-distribution-platforms/>

PRODUCT LAUNCH**Fastly**

Sep 03, 2026

Fastly Releases Open-Source Fastly Log Analytics for Self-Hosted Edge Log Insights

Fastly published Fastly Log Analytics, a self-hosted open-source data lake and interactive dashboard that lets customers ingest and query 100% of their raw Fastly request logs without a third-party analytics vendor or egress costs. It runs on Fastly Object Storage via docker-compose and includes 45 automated anomaly detectors spanning security threats, origin health, edge performance, network path, and traffic metrics, plus traffic dashboards, a global request map, and a click-to-filter raw log viewer. It is community-supported.

<https://www.fastly.com/blog/building-fastly-log-analytics-tools-open-source-approach-deeper-insights>

PARTNERSHIP**HPE / Juniper Networks**

Sep 03, 2026

HPE and Oracle Deepen Networking Collaboration for Gigawatt-Scale AI Infrastructure

HPE and Oracle expanded their networking partnership to support Oracle Cloud Infrastructure's gigawatt-scale AI buildout. The multi-year agreement deploys HPE Juniper Networking PTX and MX routing platforms plus QFX and EX switches across OCI data centers, edge networks, and AI superclusters worldwide, along with networking support services and financing. The companies are jointly developing intelligent telemetry to detect queue buildup, packet loss, and traffic imbalances that degrade GPU utilization. HPE also issued warrants allowing Oracle to purchase HPE stock.

<https://aithority.com/it-and-devops/cloud/hpe-and-oracle-deepen-networking-collaboration-to-accelerate-gigawatt-scale-ai-infrastructure/>

PARTNERSHIP**HPE / Juniper Networks**

Sep 03, 2026

HPE Inks Blockbuster Juniper Networking Deal With Oracle for AI Data Centers

CRN reported that HPE's expanded Oracle agreement is one of the industry's largest AI infrastructure networking initiatives. Oracle will deploy HPE's QFX data center switches, built on Broadcom Tomahawk silicon, alongside PTX routing platforms using HPE's own Express 5 ASIC, across AI clusters, regional data centers, and edge networks. CEO Antonio Neri framed the win as validation of HPE's post-acquisition Juniper portfolio and its AI-driven network operations story, positioning networking as the core of HPE's AI infrastructure strategy.

<https://www.crn.com/news/networking/2026/hpe-inks-blockbuster-juniper-networking-deal-with-oracle>

EARNINGS**HPE / Juniper Networks**

Sep 03, 2026

HPE Posts Record Q3 as AI Infrastructure and Networking Demand Surges

HPE reported a record fiscal Q3 2026 with revenue of \$12.2 billion, up 34% year over year, driven by AI infrastructure, networking, and servers. Campus and branch networking revenue hit record levels as customers modernize edge infrastructure and adopt AI-driven network operations, while routing and data center switching orders ran ahead of revenue and backlog reached an all-time high. HPE guided fiscal 2027 networking revenue up 13-17%. Component shortages remain a constraint, prompting multi-year supply agreements.

<https://www.networkworld.com/article/4218314/hpes-record-q3-ai-infrastructure-networking-demand-drives-growth-but-supply-constraints-tap-the-brakes.html>

PARTNERSHIP**HPE / Juniper Networks**

Sep 03, 2026

Arrow Expands HPE Networking Distribution to Aruba and Juniper Portfolios

Arrow Electronics broadened its distribution agreement with HPE to include HPE Aruba Networking and HPE Juniper Networking across the United States, Canada, and select European markets. The expansion gives partners access to HPE's full networking portfolio spanning campus, data center, cloud-managed networking, and AI infrastructure. The move is part of HPE's continued restructuring of its distribution ecosystem following the Juniper acquisition, complementing global distributors TD SYNEX and Ingram Micro and specialist distributors already in place.

<https://www.channelinsider.com/channel-business/arrow-hpe-networking-distribution/>

EXECUTIVE INTERVIEW**Barracuda Networks**

Sep 03, 2026

Barracuda CIO Details Flexible Governance Framework for Responsible AI

InformationWeek profiled how Barracuda CIO Siroui Mushegian built the company's responsible AI strategy. Barracuda formed an AI council, chaired by Mushegian, to govern internal AI use, set policy, and establish organizational tone, backed by deliberately flexible frameworks that can be revised as AI capabilities shift. The approach emphasizes visibility and access controls over how staff use data and tools to prevent shadow AI, plus cross-functional subcommittees spanning go-to-market and product engineering with a core compliance and legal group.

<https://www.informationweek.com/it-leadership/how-barracuda-s-cio-developed-a-flexible-strategy-for-responsible-ai>