


F5 Daily Briefing - Friday, September 04 - Sunday, September 06, 2026


Daily competitive intelligence covering 22 F5 competitors



Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (32)

EARNINGS Palo Alto Networks 2026-09-06

Palo Alto Networks Stock Falls 10% Despite Strong Q4 FY26 Results and Raised FY27 Guidance

Palo Alto Networks shares dropped roughly 10% during the week following Q4 fiscal 2026 results. Revenue rose about 34% year over year to \$3.41 billion with adjusted EPS of \$1.02, next-generation security ARR climbed 63% to \$9.1 billion, and remaining performance obligations grew 34% to \$21.2 billion. FY2027 guidance of \$14.1-14.2 billion in revenue came in above consensus. The pullback reflected elevated expectations after an approximately 81% year-to-date run rather than operational weakness.

Source: <https://www.fool.com/investing/2026/09/06/why-palo-alto-networks-stock-plummeted-this-week/>

RESEARCH Palo Alto Networks 2026-09-05

Unit 42 Documents AI Agents Breaching an Enterprise Network in Under 10 Hours

Palo Alto Networks' Unit 42 published research on an intrusion where attackers used frontier AI models and agentic frameworks to compromise an enterprise network in under 10 hours. The agents mapped internal services, searched code repositories for secrets, reached a secrets-management platform, and obtained root credentials, chaining more than 50 MITRE ATT&CK techniques and repurposing the victim's own AI endpoints and cloud compute. The report frames machine-speed attacks as compressing multiple intrusion phases into a single automated loop.

Source: <https://cybersecuritynews.com/ai-agents-breach-company-network/>

RESEARCH Lasso Security 2026-09-05

Lasso's CPU-Only LEAP Guardrail Draws Scrutiny Over In-House Benchmarks

Follow-on coverage of Lasso Security's LEAP guardrail, a transformer-free engine that inspects AI requests and agent actions on ordinary CPUs with claimed sub-5ms verdicts and no GPU footprint, flagged that the performance figures are vendor-run and not independently validated. Reporting notes Lasso has not disclosed CPU model, core allocation, test datasets, concurrency or measurement methodology, and offers no reproducible detection-quality comparison against rival guardrails. The CPU fast path and Lasso's \$30M ClearSky-led round are confirmed; the throughput advantage remains unverified.

Source: <https://quasa.io/insights/lasso-s-guardrail-runs-on-cpus-in-under-5ms-its-benchmark-is-in-house>

EARNINGS **Fastly** 2026-09-05

Fastly Lifts 2026 Guidance as AI Infrastructure Demand Drives Revenue Beat

Fastly raised its full-year 2026 guidance after beating expectations, crediting stronger demand for AI infrastructure and edge cloud services. Revenue grew 23.3% year over year to \$183.3 million, with Security up 43% to \$41.7 million and Compute and Observability up 69% to \$7.7 million on AI-related traffic. Management cited AI-generated and agentic traffic as a key demand catalyst and highlighted sub-millisecond Compute cold-start latency against hyperscaler and CDN rivals. The top ten customers represent roughly 38% of revenue.

Source: <https://www.streamingmeme.com/articles/fastly-lifts-2026-guidance-as-ai-infrastructure-demand-drives-revenue-beat>

PRODUCT LAUNCH **CrowdStrike** 2026-09-05

Falcon Guardian Makes the Endpoint the Enforcement Layer for AI Agents, With AI Gateway Coming in Q4

CrowdStrike positioned Falcon Guardian as a runtime enforcement layer that turns the endpoint into the primary control point for AI agent governance across Windows and macOS. It discovers known and shadow agents and links prompts to tool calls and downstream actions, with claims of 99% prompt-attack detection at roughly 100ms latency. CrowdStrike also signaled a new AI Gateway planned for Q4 2026 to centralize enterprise AI traffic, moving beyond visibility and posture management into real-time enforcement.

Source: <https://forkast.news/crowdstrike-falcon-guardian-makes-the-endpoint-the-enforcement-layer-for-ai-agents/>

SECURITY ADVISORY **CrowdStrike** 2026-09-05

FalconFlank Zero-Day Turns CrowdStrike Falcon Sensor Into a Privilege Escalation Path

A researcher published FalconFlank, a working privilege-escalation exploit for CrowdStrike Falcon Sensor on Windows that abuses the product's own Office malicious-macro removal routine to spawn a SYSTEM-level shell. Affected configurations include Windows 11 25H2 and Windows Server 2025 with the Suspicious Macro Removal policy enabled. CrowdStrike confirmed an investigation, published a customer Tech Alert, and advised disabling that policy setting as an interim mitigation. No CVE or patch had been issued and no in-the-wild exploitation was confirmed.

Source: <https://www.techtimes.com/articles/326752/20260905/crowdstrike-falcon-zero-day-turns-enterprise-security-agent-attack-surface.htm>

COMPANY NEWS **Cisco** 2026-09-05

Cisco Scans 1.8 Billion Lines of Code as AI Accelerates Vulnerability Discovery

Cisco detailed how AI-assisted analysis across roughly 1.8 billion lines of its own code is accelerating internal vulnerability discovery, the program behind its recent proactive security hardening releases. Executives also signaled a target of having effectively no engineers hand-writing code by late October as AI tooling takes over more of the software supply chain. The move positions AI-driven secure development as a differentiator in Cisco's platform narrative and reshapes expectations for vendor patch cadence.

Source: <https://insight.tmcnet.com/insight/cisco-scans-1-8-billion-lines-as-ai-accelerates-vulnerability-discovery-mto6derz>

COMPANY NEWS **AWS** 2026-09-05

AWS Commits \$5.3B and 50MW of AI Capacity to Saudi AI Zone with HUMAIN

At the LEAP conference in Riyadh, AWS reaffirmed that its first Saudi Arabia cloud region remains on track for December 2026, expanding its footprint to 40 regions. AWS also disclosed a commitment to supply up to 50 megawatts of AI computing capacity through a shared AI Zone facility with HUMAIN, the Saudi state-backed AI

company, by 2028. The \$5.3 billion investment ties into broader Saudi AI infrastructure and workforce programs targeting more than 100,000 citizens.

Source: <https://ppc.land/aws-commits-5-3-billion-50mw-capacity-to-saudi-ai-zone-by-2028/>

EARNINGS Zscaler 2026-09-04

Zscaler Beats on Q4 FY2026 but Guides FY2027 Sharply Lower and Cuts 3% of Workforce

Zscaler reported Q4 FY2026 revenue of \$898.2 million, up 25% year over year, with adjusted EPS of \$1.19 versus \$1.09 expected. FY2027 guidance calls for revenue and ARR growth of just 16.6%-17.5%, a marked deceleration. The company paired the outlook with a roughly 3% global headcount reduction and \$30-\$33 million in restructuring charges, mostly in the first half of FY2027. Free cash flow margin fell to 6.8% from 16% on accelerated data-center hardware purchases, and shares dropped pre-market.

Source: <https://tradevae.com/news/stock-markets/zscaler-shares-fall-after-strong-quarter-but-weaker-forward-growth-forecast/>

PRODUCT LAUNCH Zscaler 2026-09-04

Zscaler Unveils OpenAI-Built Endpoint Attack Chain and Identity Risk Analysis Capabilities

Zscaler introduced two AI-powered security capabilities developed with OpenAI using GPT cyber models. Endpoint AI Security Attack Chain Analysis reasons over full endpoint state, including AI agents, IDEs, and browser extensions, to reconstruct realistic multi-step attack chains instead of listing isolated vulnerabilities. Identity-Permission Risk Analysis plugs into Zscaler's AI Access Graph to evaluate humans, service accounts, and AI agents across Azure, AWS, OCI, and GCP, ranking risky permission patterns by business impact. Both are evidence-grounded and require human review before consequential action.

Source: <https://insidecybersec.com/zscaler-unveils-ai-powered-security-capabilities-built-with-openai/>

PRODUCT LAUNCH Zscaler 2026-09-04

Zscaler Demonstrates ZPA Paired With AI Guard for Private AI Deployments

Zscaler published a walkthrough of Zscaler Private Access working with AI Guard to secure private AI and ML deployments. ZPA segments and tags sanctioned AI applications so policy guardrails apply based on business context, user group, and risk, while AI Guard blocks prompts containing restricted data such as PII or prohibited formats in real time. ZPA diagnostics and the AI Guard dashboard show which policy fired and why. The post positions inline prompt inspection as a native extension of zero trust private access.

Source: <https://www.zscaler.com/blogs/product-insights/zpa-and-ai-guard-action>

RESEARCH Radware 2026-09-04

Radware Research Warns Attackers Are Hijacking Trusted AI Assistants for Phishing

Radware published new research arguing that phishing is shifting away from deceptive emails toward manipulating the AI assistants users already trust. Via prompt injection, attackers can poison the answers an assistant returns, turning a trusted interface into the delivery channel for malicious instructions and links. Radware ties the findings to its broader work on indirect prompt injection and AI agent protection, positioning prompt-injection defense and agent-activity visibility as core requirements rather than optional add-ons for enterprises deploying AI assistants.

Source: https://www.linkedin.com/posts/radware_aiphishing-cybersecurity-promptinjection-activity-7501609919263981568-Mo1C

COMPANY NEWS Palo Alto Networks 2026-09-04

Analysts Frame Console Acquisition as Palo Alto's Move From AI Copilots to Agentic SecOps

PAC analysts assessed Palo Alto Networks' acquisition of Console, arguing it marks a shift from AI copilots that suggest actions to agentic security operations that execute them. Console folds into Cortex, the security operations platform Palo Alto has rebuilt through successive acquisitions, and is positioned to deliver autonomous alert resolution and natural-language remediation workflows. The analysis frames platform consolidation plus agentic execution as the next competitive battleground in security automation.

Source: <https://sitsi.pacanalyst.com/from-ai-copilots-to-agentic-secops-what-palo-alto-networks-acquisition-of-console-tells-us-about-the-next-phase-of-security-automation/>

FUNDING **Mindgard** 2026-09-04

Mindgard Closes \$30M Series A to Scale Automated AI Red Teaming

Lancaster University spin-out Mindgard confirmed a US\$30 million (22.2M pound) Series A led by Album VC with Karma Ventures and existing backers .406 Ventures, Atlantic Bridge, IQ Capital and Lakestar, one of the largest UK cybersecurity Series A rounds. Funds go to product engineering, sales and marketing for its continuous automated red-teaming platform, which has surfaced 150+ high-impact AI vulnerabilities including a Cursor IDE zero-day and flaws in Google Antigravity and ChatGPT. Mindgard operates from London and Boston across finance, pharma, gaming and semiconductors.

Source: <https://www.lancaster.ac.uk/news/major-funding-boost-for-lancaster-ai-spin-out-company>

PARTNERSHIP **HPE / Juniper Networks** 2026-09-04

Oracle Expands HPE Juniper Networking Deal to Power Gigawatt-Scale AI Infrastructure

Oracle significantly expanded its multi-year agreement with HPE Juniper Networking to scale Oracle Cloud Infrastructure's AI data centers, edge networks, and regional hubs. The deployment covers PTX and MX routing platforms plus high-density QFX and EX switches supporting RoCEv2 AI back-end fabrics for large GPU superclusters, with congestion management, fault tolerance, and intelligent telemetry to protect GPU utilization. The deal bundles multi-year networking support services and financing, and includes warrants letting Oracle purchase HPE common stock.

Source: <https://www.itpro.com/infrastructure/networking/oracle-expands-hpe-juniper-networking-deal-to-support-ai-infrastructure-build-out>

COMPANY NEWS **HPE / Juniper Networks** 2026-09-04

Tottenham Hotspur Replaces VMware With HPE Stack, Cutting Licensing Costs 85%

Tottenham Hotspur migrated its stadium IT off VMware to HPE Morpheus VM Essentials, cutting virtualization licensing costs by more than 85% while consolidating eight environments into two core sites. The refresh spans ProLiant Gen12 servers, Alletra Storage MP, Aruba Central for hybrid cloud network management, and OpsRamp for multi-site operations, supporting IPTV, POS, crowd analytics, and Wi-Fi 7 across 1,650-plus access points. HPE positions the build as an AI-native networking backbone; Juniper integration with Aruba/Mist is expected in later phases.

Source: <https://www.techradar.com/pro/tottenham-hotspur-ditches-vmware-in-favor-of-hpe-cuts-licensing-costs-by-85-percent>

COMPANY NEWS **Fortinet** 2026-09-04

Fortinet Joins Project Watershed 250 to Harden National Water Cybersecurity Infrastructure

Fortinet announced participation in Project Watershed 250, a six-month pilot uniting federal and Texas state agencies, water utilities, and private cybersecurity firms to strengthen the resilience of Texas water infrastructure. Fortinet is contributing offense/defense testing, hardening support, threat intelligence, and AI-enabled defenses. Led by the White House National Cyber Director alongside Texas Cyber Command, CISA, and EPA, the program aims to identify and remediate vulnerabilities and share lessons across the wider water sector.

Source: <https://www.fortinet.com/blog/industry-trends/fortinet-joins-project-watershed-250-strengthen-national-water-cybersecurity-infrastructure>

PRODUCT LAUNCH **Fastly** 2026-09-04

Fastly Releases Open-Source Self-Hosted Log Analytics Data Lake and Dashboard

Fastly announced Fastly Log Analytics, an open-source, self-hosted data lake and interactive dashboard that captures 100% of raw request-level Fastly logs and runs on the customer's own infrastructure using Fastly Object Storage. The stack uses Apache Iceberg tables and DuckDB for fast local querying, eliminating third-party observability vendor bills. Fastly positions it as delivering request-level visibility beyond aggregate dashboards, letting teams identify exactly who is behind traffic spikes or widespread 403 responses.

Source: https://www.linkedin.com/posts/fastly_deeper-log-visibility-without-the-third-party-activity-7501694358040236032-9gU7

SECURITY ADVISORY Fastly 2026-09-04

Fastly Ships Next-Gen WAF Virtual Patch as JFrog Artifactory Exploitation Escalates

Fastly's Threat Research team reported a sharp escalation in exploitation attempts against CVE-2026-82329, a critical authentication bypass in self-hosted JFrog Artifactory. Since disclosure on August 28, attacker behavior shifted from targeted probing to opportunistic mass scanning. Fastly said its Next-Gen WAF has shipped a virtual patch that protects affected instances immediately, giving teams unable to upgrade Artifactory a mitigation path. Full technical analysis and remediation guidance were published in an accompanying Fastly blog post.

Source: https://www.linkedin.com/posts/fastly_jfrog-artifactory-waf-activity-7501437601422438400-YzqY

PRODUCT LAUNCH CrowdStrike 2026-09-04

CrowdStrike Unveils Falcon Guardian AIDR and SafeMind Agentic Defense at Fal.Con 2026

At Fal.Con 2026, CrowdStrike launched Falcon Guardian, an AI detection and response (AIDR) platform that discovers and maps known and shadow AI agents across endpoints, then governs and contains AI-driven threats in real time. It also introduced SafeMind, a defender-focused agentic AI system built with NVIDIA Nemotron that uses frontier models and a digital-twin red-team/blue-team environment to harden defenses. CrowdStrike frames both as securing the AI revolution at the point where agents execute.

Source: <https://biztechmagazine.com/article/2026/09/crowdstrike-falcon-2026-crowdstrike-announces-falcon-guardian-and-safe-mind-aidr>

PARTNERSHIP CrowdStrike 2026-09-04

VAST Data and CrowdStrike Integrate Security Across the AI Operating System

CrowdStrike and VAST Data announced an integration that extends enterprise security across infrastructure, data, and AI workloads in production AI environments. The native CrowdStrike Falcon sensor now supports VAST, with Falcon Next-Gen SIEM and Falcon Guardian AIDR in private preview on the VAST AI Operating System. The integration pushes protection into AI pipelines, data access, and data preparation for anomaly detection, and will be demonstrated on an NVIDIA-powered AI factory alongside VAST's accelerated data pipelines.

Source: <https://investingnews.com/vast-data-and-crowdstrike-deliver-first-of-its-kind-cybersecurity-across-the-ai-operating-system/>

PRODUCT LAUNCH Cloudflare 2026-09-04

Cloudflare Raises Workers Bundle Limit to 64 MiB on Free and Paid Plans

Cloudflare removed the compressed-bundle deployment ceiling for Workers, previously 3 MiB on Free and 10 MiB on Paid plans, and now enforces only a 64 MiB uncompressed bundle size across both tiers. Developers can verify size before shipping using `wrangler deploy --outdir bundled/ --dry-run`, where Total Upload reflects the uncompressed bundle. The change lets substantially larger applications, dependencies, and AI/agent workloads run directly on Cloudflare's edge compute platform without bundle-splitting workarounds.

Source: <https://developers.cloudflare.com/changelog/post/2026-09-04-increased-worker-size-limit/>

SECURITY ADVISORY Citrix (NetScaler) 2026-09-04

NetScaler ADC Authentication Bypass CVE-2026-19490 Now Exploited in the Wild

Australia's ACSC and the Canadian Centre for Cyber Security issued urgent advisories on critical NetScaler ADC and NetScaler Gateway flaws CVE-2026-19490 (unauthenticated authentication bypass, CVSS 9.3) and CVE-2026-19489 (memory overflow via SIP ALG). Following a public proof-of-concept, sensors observed exploitation attempts beginning September 3 from IPs in Australia, the US, and Germany, and the bug was added to the known-exploited catalog. Shadowserver counted more than 22,000 exposed ADC and roughly 1,700 Gateway instances. Fixed builds are 14.1-73.32 and 13.1-63.21 or later.

Source: <https://australiancybersecuritymagazine.com.au/acsc-warns-of-critical-vulnerabilities-in-citrix-netScaler-adc-and-netScaler-gateway/>

SECURITY ADVISORY Cisco 2026-09-04

Cisco Patches Critical Nexus 9000 RCE and IOS XR Flaws in Internal Security Review

Cisco disclosed and fixed three critical vulnerabilities found during a proactive internal security review, including CVE-2026-20212 (CVSS 9.8) in Nexus 9000 Series switches using Silicon One ASICs. The flaw allows

unauthenticated remote code execution as root via TCP ports 43210 and 43211 exposed in the default Layer 3 VRF, and can crash the S1HAL process. Ten Nexus 9000 models are affected, with no workarounds beyond access control lists; Cisco also shipped an IOS XR hardening release.

Source: <https://www.techradar.com/pro/security/cisco-patches-three-critical-vulnerabilities-as-part-of-comprehensive-internal-security-review>

PRODUCT LAUNCH Cisco 2026-09-04

Cisco Adds Agent Runtime Protection to AI Defense Python SDK

Cisco released Agent Runtime Protection in the Cisco AI Defense Python SDK, giving developers a one-line integration that dynamically rewrites LLM and MCP client calls to route through AI Defense inspection without changing application code. The capability inspects prompts, responses, and MCP interactions including tools, prompts, and resources, screening for prompt injection and data exfiltration before and after each call. It targets teams operationalizing agentic AI who need runtime guardrails embedded in the agent framework itself.

Source: <https://dailysynapse.com/news/cisco-adds-runtime-protection-for-ai-agents/>

PARTNERSHIP Check Point Software 2026-09-04

Check Point Embeds OpenAI Daybreak Frontier Models Across Its Security Platform

Check Point expanded its OpenAI Daybreak collaboration, embedding frontier cyber reasoning models across its platform through the Daybreak Defense Network. Integrations span Agentic Exposure Validation in exposure management, Keystone agentic security management, the Autonomous Workspace Platform for cross-telemetry incident investigation, and NexPloit vulnerability research. A multi-agent pipeline separates genuinely exploitable risk from theoretical findings and validates attack paths using Check Point context and evidence. Rollout is phased, with governance controls limiting what models can see and do, and output verification before any action.

Source: <https://securitymea.com/2026/09/04/check-point-brings-openai-daybreak-models-across-its-platform/>

COMPANY NEWS Barracuda Networks 2026-09-04

Barracuda Managed XDR Named a Top Profit-Driving MSP Solution for 2026

Barracuda announced that Barracuda Managed XDR was recognized by ChannelPro among its 26 Profit-Driving Solutions for MSPs in 2026. The company frames Managed XDR as a way for partners to deliver 24/7 threat monitoring, detection, investigation, and response without the cost of building an in-house SOC. Barracuda emphasizes recurring-revenue growth, reduced alert fatigue, and deeper customer engagement around risk, compliance, and business continuity, reinforcing its partner-first, MSP-led platform motion built around BarracudaONE.

Source: <https://blog.barracuda.com/2026/09/04/barracuda-managed-xdr-top-msp-solution-2026>

EARNINGS Arista Networks 2026-09-04

Arista's AI Networking Momentum Draws Fresh Investor Scrutiny After Q2 Beat

Arista Networks drew renewed analyst attention over its AI, cloud, and enterprise networking momentum, having posted roughly \$3.036 billion in quarterly revenue and \$1.213 billion in net income, both up year over year, with guidance near \$3.3 billion for the next quarter. Coverage emphasized Arista's position in high-speed AI infrastructure Ethernet and its EOS and CloudVision software stack, alongside diversification into enterprise campus and branch networking. Arista is also slated to present at the Goldman Sachs Communacopia and Citi Global TMT conferences in September.

Source: <https://sg.finance.yahoo.com/news/arista-q2-ai-networking-beat-131003114.html>

PRODUCT LAUNCH AWS 2026-09-04

AWS MCP Server Adds Serverless Diagnostics for Lambda Functions

AWS extended its MCP Server with a serverless capability that lets coding agents such as Claude Code and Kiro diagnose Lambda functions and connected resources including API Gateway, EventBridge, S3, DynamoDB, SNS, SQS and Step Functions. It correlates errors against a seven-day baseline, surfaces recurring issues, retrieves deployed configurations, tracks recent changes and analyzes service latency, returning results in a single call to reduce token usage. Available via the Agent Toolkit for AWS at no additional cost.

PRODUCT LAUNCH **AWS** 2026-09-04

Amazon Bedrock Managed Knowledge Base Adds Automatic Sync Scheduling

Amazon Bedrock Managed Knowledge Base now supports automatic sync scheduling for data source connectors, letting customers keep retrieval-augmented generation indexes current without manually triggering ingestion jobs or building custom orchestration. The capability reduces operational overhead for teams running production RAG and agentic applications on Bedrock, where stale knowledge bases degrade answer quality. It reinforces AWS's push to make managed AI data pipelines a default part of its inference stack rather than customer-built glue code.

Source: <https://aws.amazon.com/about-aws/whats-new/2026/09/amazon-bedrock-managed-knowledge-base-automatic-sync-scheduling/>

PRODUCT LAUNCH **AWS** 2026-09-04

AWS Details Open-Source HyperPod InstantStart Control Plane for Agent Operations

AWS published details on HyperPod InstantStart, an open-source control plane that pairs Amazon EKS orchestration with SageMaker HyperPod's managed training and inference capabilities. It ships with a web interface plus an AI agent that plans and executes multi-stage cluster operations using Model Context Protocol tools. The release targets teams standing up large GPU clusters for AI workloads, lowering the setup and day-two operations burden for distributed model training and inference infrastructure.

Source: <https://tradepoint.io/aws-details-open-source-hyperpod-instantstart-control-plane-for-agent-ops-unite-ai/>

COMPANY NEWS **A10 Networks** 2026-09-04

A10 Networks Shares Pull Back as AI Infrastructure Narrative Keeps Valuation in Focus

A10 Networks stock closed at \$24.71 on September 3 and is down roughly 19% over the past month, despite remaining up about 42% year to date. Analyst commentary frames the pullback against A10's AI infrastructure positioning, pointing to data center expansion and demand for its secure application and network portfolio as the core growth thesis, with a fair value estimate of \$38.40. The narrative rests heavily on A10 converting its recent AI Gateway and TrojAI AI security additions into durable ADC and security revenue.

Source: <https://simplywall.st/stocks/us/software/nyse-aten/a10-networks/news/a10-networks-aten-pulls-back-as-the-ai-infrastructure-narrat>