

F5 Daily Briefing - Monday, September 21, 2026

Daily competitive intelligence covering 22 F5 competitors

22

COMPETITORS TRACKED

19

ANNOUNCEMENTS


11

ACTIVE


3

ACTIVE CATEGORIES

Announcements by Competitor



Announcements by Category



ALL ANNOUNCEMENTS (19)

FINANCIAL

Arista Networks

2026-09-21

Arista Networks Joins S&P 100 as AI Infrastructure Weighting Grows

Arista Networks joined the S&P 100 index effective before market open on September 21, 2026, replacing Nike in S&P Dow Jones Indices' quarterly rebalancing alongside moves involving Dell, Palo Alto Networks and SanDisk that tilt the index further toward AI infrastructure names. Shares rose roughly 1.5% on the day to about \$202.56. Analyst sentiment remains strongly positive, with 120 of 122 covering analysts rating the stock a buy and a median price target near \$246.50. The inclusion reflects investor recognition of Arista's AI data center networking growth.

<https://tradersunion.com/news/companies/show/3418107-arista-networks-stock-rises-1-5percent/>

PRODUCT LAUNCH

AWS

2026-09-21

AWS Pattern Streams Bedrock Guardrails Interventions to OCSF in CloudWatch Unified Store

This AWS architecture pattern normalizes Amazon Bedrock guardrail interventions — harmful content blocks, PII leaks, and prompt-injection attempts — into the Open Cybersecurity Schema Framework DetectionFinding format and centralizes them in the CloudWatch unified data store. A CDK-deployed Lambda subscribes to Bedrock model invocation logs, maps guardrail records to OCSF, and writes JSON that can be queried alongside IAM, VPC Flow Logs, and CloudTrail data via Athena or CloudWatch Logs Insights. The result is cross-service correlation of AI guardrail alerts with other security telemetry, complementing GuardDuty AI Protection findings.

<https://cloudninja.ca/ai/streaming-bedrock-guardrails-to-ocsf-in-the-cloudwatch-unified-store/>

PRODUCT LAUNCH**AWS**

2026-09-21

AWS Releases Strands Harness, an Open-Source Multi-Cloud AI Agent Harness

AWS debuted Strands Harness, an open-source, general-purpose agent harness built on the Strands Harness SDK that deploys in any environment including AWS, Google Cloud, Azure, Modal, and Cloudflare. It ships Python and TypeScript libraries, a CLI, and a skills file with shell, file, and web tools, and works with any AI model provider. AWS claims near-parity benchmark scores with Claude Code and Codex on Harbor tests while using about 28% fewer tokens, attributing the efficiency to default prompt caching, truncation of large tool results, auto-compaction at 85% context usage, and context recovery on overflow.

<https://siliconangle.com/2026/09/21/aws-debuts-strands-harness-an-open-source-ai-agent-that-can-be-deployed-in-any-environment/>

PARTNERSHIP**AWS**

2026-09-21

Amazon Bedrock Adds Kimi K3 Under a New Model Revenue-Sharing Partnership

Amazon Bedrock integrated the open-source large language model Kimi K3 into the AWS AI service platform, giving global enterprise developers direct access to the model through Bedrock. The integration marks the formal launch of Kimi's revenue-sharing partnership with international cloud providers, extending Bedrock's third-party model catalog and reinforcing AWS's strategy of positioning Bedrock as the aggregation point for frontier and open-weight models that enterprises consume through a single managed inference layer.

<https://www.gate.com/news/detail/amazon-bedrock-integrates-kimi-k3-launching-revenue-sharing-model-with-24439709>

THREAT INTEL**Check Point Software**

2026-09-21

Check Point Patches CVSS 9.8 Unauthenticated Root RCE in Security Management and Log Servers

Check Point issued an emergency patch for CVE-2026-91843, a CVSS 9.8 stack-based buffer overflow in the Security Management Server and Log Server login process (R80-R82) that gives unauthenticated attackers root access over the internet. It is the third critical Check Point flaw disclosed in two weeks; two earlier bugs, CVE-2026-50751 and CVE-2026-16232, are already exploited in the wild, and CVE-2026-85103 and CVE-2026-85102 were patched the prior week. Check Point issued patches plus interim guidance to restrict management console access to trusted IP addresses until customers can upgrade.

<https://threatvectr.com/story/check-point-patches-critical-login-flaw-that-hands-attackers-root-on-firewall>

RESEARCH**Check Point Software**

2026-09-21

Check Point Research Says Agent Security Must Watch Behavior, Not Just Attackers

Check Point published research arguing AI agents can cause harm without any attacker present, simply by acting on their context. Across a two-day internal hackathon with 19 R&D teams, three independent demonstrations converged on the same conclusions: agents improvise when they hit walls, every file, page, and message in an agent's context is an attack surface, and blocking is not the only guardrail — steering an agent back to its intended task preserves security with less productivity loss. Check Point paired the post with a white paper covering context poisoning and runtime protection for production agents.

<https://blog.checkpoint.com/ai-security/no-attacker-required-what-a-two-day-hackathon-taught-us-about-agent-security/>

RESEARCH**Cisco**

2026-09-21

Cisco 2026 Global AI Partner Study Finds AI Security Awareness Falling as Advisory Demand Surges

Cisco released its 2026 Global AI Partner Study, surveying 1,702 partners across 30 countries. Partner awareness of AI-specific security threats fell to 29% from 35% a year earlier, while security and compliance ranked as partners' second-largest internal skills gap and cybersecurity emerged as a top customer AI talent gap. Governance and explainability remain weak: partners can identify AI risks but struggle to explain model decisions or demonstrate regulatory compliance. Advisory and consulting services grew from 12% to 50% of partner opportunity. Cisco is responding with an AI-Era Cyber Defense Partner Hub and its 360 Partner Program.

<https://blogs.cisco.com/partner/partner-study-shows-a-maturing-ai-market-expanding-opportunity-for-partners>

PRODUCT LAUNCH**Cisco**

2026-09-21

Cisco Extends MACsec Encryption to 800ZR Coherent Optics for Secure AI Networks

Cisco detailed a silicon-to-optics platform strategy for the agentic AI era, arguing AI agents can drive up to 450% more network traffic than traditional workflows, with roughly 70% coming from AI inference. Cisco says securing that traffic matters as much as transporting it, and is extending MACsec (IEEE 802.1AE) link encryption to the optical layer via 800ZR coherent pluggables so operators can secure high-capacity optical transport without adding separate encryption devices. Cisco also highlighted quantum-resistant cryptographic readiness and is demonstrating the optical, routing, and secure transport portfolio at ECOC 2026.

<https://blogs.cisco.com/news/the-ai-era-demands-more-than-speed-building-secure-intelligent-networks-from-silicon-to-optics>

THREAT INTEL**Cisco**

2026-09-21

Cisco Identity Services Engine Authentication Bypass CVE-2026-76460 Under Active Exploitation

Singapore's Cyber Security Agency and multiple advisories flagged active in-the-wild exploitation of CVE-2026-76460, a maximum-severity authentication bypass (CVSS 10.0) in Cisco Identity Services Engine and ISE Passive Identity Connector. Successful exploitation can grant highly privileged root-level access and remote command execution, and the flaw affects all configurations with no available workaround. Cisco has released fixed releases and urges immediate upgrade; organizations still running ISE 3.0, which has reached end of maintenance, must migrate to a supported fixed train. Defenders are advised to verify affected versions against Cisco's advisory urgently.

<https://www.csa.gov.sg/alerts-and-advisories/alerts/al-2026-126/>

RESEARCH**Cisco**

2026-09-21

Cisco Posts Industry-First Multi-Vendor MLPerf Inference Results Across NVIDIA and AMD Accelerators

Cisco published what it calls the industry's first MLPerf Inference submission spanning multiple accelerator vendors, running distributed inference across eight NVIDIA H200 GPUs and eight AMD Instinct MI350X GPUs connected over a Cisco Silicon One G200 Ethernet fabric for MLPerf Inference v6.1. Using vLLM-based distributed inference, Cisco reported up to 104.7% of the combined standalone throughput of the two pools, indicating no net penalty for cross-architecture operation. Cisco positions the result as proof that standard Ethernet fabrics can pool heterogeneous GPUs, improving inference unit economics and freeing enterprises from single-vendor accelerator lock-in.

<https://blogs.cisco.com/datacenter/industrys-first-mlperf-inference-benchmarking-across-multi-vendor-accelerators>

PRODUCT LAUNCH**Cloudflare**

2026-09-21

Cloudflare Adds Per-Worker Scoped Access Roles from the Dashboard

A Cloudflare changelog entry added the ability to grant teammates scoped access to individual Workers directly from the Workers dashboard. Administrators open a Worker, choose Invite, enter an email address, and select one of four access levels: Metadata Read-Only (settings, metrics, logs, traces but no code or deploy), Content Read-Only (read code, settings, and observability data with no modify or deploy), Editor (update and deploy but not delete), or Admin (Editor plus delete). Existing account members gain access immediately; others receive an invitation. Only Super Administrators can issue invites from a Worker's dashboard.

<https://developers.cloudflare.com/changelog/post/2026-09-21-invite-members-to-workers/>

PRODUCT LAUNCH**Cloudflare**

2026-09-21

Cloudflare Makes Python Workers Generally Available on Its Developer Platform

Cloudflare announced that Python Workers are generally available, making Python a first-class, fully supported language on the Cloudflare Developer Platform two years after its initial preview. Developers can bring existing Python code and libraries and run popular frameworks including FastAPI, Django, and Flask directly inside Workers. Python Workers support native bindings to Workers AI, R2, D1, Hyperdrive, Durable Objects, Queues, and Workflows, and a Python Worker can be created from inside another Worker using Dynamic Workers. The release removes the prior requirement for JavaScript glue code when running Python at Cloudflare's edge.

<https://blog.cloudflare.com/python-workers-ga/>

PRODUCT LAUNCH**Cloudflare**

2026-09-21

Cloudflare Proposes an Agent Development Lifecycle to Replace the Traditional SDLC

Cloudflare introduced the Agent Development Lifecycle (ADLC), a model in which AI agents autonomously handle development, testing, deployment, and maintenance instead of human-gated pipelines. Components include Workflows as the orchestration layer that spawns containers and dispatches subagents, @cloudflare/ci as a CI/CD system running on Workflows with autonomous failure triage and remediation, OpenTelemetry-compatible observability capturing model calls, tool executions and token usage with session replay, and an Agent Access Model issuing short-lived task-bound credentials with least-privilege enforcement and a Trust Ratchet that progressively reduces agent capabilities.

<https://www.infoq.com/news/2026/09/cloudflare-adlc-agents/>

PARTNERSHIP**CrowdStrike**

2026-09-21

Horizon3 Integrates NodeZero Exposure Validation with CrowdStrike Falcon Next-Gen SIEM

Horizon3 announced an integration that feeds validated NodeZero findings into CrowdStrike Falcon Next-Gen SIEM, now available in the CrowdStrike Marketplace. Validated exposure data showing genuinely exploitable attack paths can be correlated with endpoint, identity, and cloud telemetry inside Falcon investigations, helping analysts prioritize which findings matter and tune policy accordingly. The integration closes the loop between proactive offensive validation and runtime protection, reinforcing CrowdStrike's strategy of pulling third-party security signals into Falcon Next-Gen SIEM as the central data and correlation layer rather than leaving validation results in a separate console.

<https://www.securitysolutionsmedia.com/2026/09/21/horizon3-announces-integration-with-crowdstrike-falcon-next-gen-siem/>

PRODUCT LAUNCH**Fastly**

2026-09-21

Fastly Launches AI Firewall and AI Runtime Control to Secure and Scale AI

Fastly expanded its "Fastly for AI" portfolio with three capabilities, all generally available now. AI Runtime Control acts as a control plane that routes model calls across public and self-hosted providers, with real-time visibility into token spend, rate limits, budgets, virtual keys for credential protection, and failover. AI Firewall inspects prompts in the request path at the edge to block prompt injection and other LLM attacks, and can be targeted at specific models. Expanded API Security enforces API contracts across agent-based, agent-assisted, and traditional traffic. Fastly cited machine-generated traffic overtaking human traffic in July 2026 as the driver.

<https://www.fastly.com/press/press-releases/fastly-launches-ai-firewall-and-ai-runtime-control-to-secure-and-scale-ai>

THREAT INTEL**Fortinet**

2026-09-21

Fortinet FortiCloud SSO Flaws Still Exploited, Leaving 10,000+ Devices Exposed

A report warned that critical FortiCloud SSO vulnerabilities CVE-2025-59718 and CVE-2025-59719 (CVSS 9.8) remain actively exploited, leaving more than 10,000 devices worldwide at risk. Attackers can bypass authentication using crafted SAML messages to obtain full administrative access, create rogue accounts, and exfiltrate firewall configurations. Patches through FortiOS 7.4.10 proved insufficient, with comprehensive fixes pending in FortiOS 7.4.11, 7.6.6, and 8.0.0. Affected products span FortiOS, FortiProxy, FortiSwitchManager, and FortiWeb. Recommended interim mitigations include disabling FortiCloud SSO login, rotating admin accounts, and auditing configurations for unauthorized changes.

<https://www.ctrlaltnod.com/news/fortinet-sso-exploit-over-10-000-devices-at-risk-despite-patches/>

RESEARCH**HPE / Juniper Networks**

2026-09-21

HPE Networking Pushes Self-Driving Networks From Vision Toward Autonomous Remediation

A BizTech Magazine feature details how HPE Networking is advancing from AI-assisted troubleshooting to genuinely autonomous, self-driving network operations. HPE combines AI, network telemetry, cloud-native microservices and years of accumulated customer trouble-ticket data to identify recurring problems and autonomously apply fixes with reduced manual intervention. HPE's Rich Carter describes AI auto-correcting common network issues as a core element of the broader self-driving network strategy, noting the approach requires extensive longitudinal data, deep understanding of network behavior, and organizational confidence to let AI take remediation action on production infrastructure.

<https://biztechmagazine.com/article/2026/09/ai-self-driving-networks-move-vision-reality>

FINANCIAL**Palo Alto Networks**

2026-09-21

Morgan Stanley Lifts Palo Alto Networks Target to \$410, Citing AI Security Demand and Consolidation

Morgan Stanley raised its Palo Alto Networks price target to \$410 from \$394, keeping an Overweight rating and naming PANW its preferred cybersecurity pick. The thesis is that enterprise AI adoption will expand security spending while pushing customers to consolidate onto fewer vendors, with AI security products, identity management, and SIEM as the main growth drivers toward a roughly \$300 billion cybersecurity software market by 2028. Morgan Stanley expects Palo Alto to gain about 75 basis points of share over two years versus consensus expectations of a 50 basis point decline.

<https://finviz.com/news/393715/morgan-stanley-raises-palo-alto-networks-price-target-to-410-on-ai-security-demand>

FINANCIAL**Zscaler**

2026-09-21

Zscaler AI Bookings Jump More Than 50% Sequentially as Security for AI Scales

Coverage of Zscaler's Q4 FY2026 results highlighted Security for AI bookings growing more than 50% sequentially, with AI Protect bookings surpassing \$100 million over the trailing twelve months. Roughly 70% of Security for AI deals also attach Data Security. Q4 revenue was \$898.2 million, up 25% year over year, and ARR reached \$3.77 billion, also up 25%. Newer products include Zero Trust Exchange for Agents and Endpoint AI Security, with broader scaling expected in the second half of FY2027. FY2027 guidance implies roughly 17% revenue growth, materially slower than FY2026.

<https://sg.finance.yahoo.com/news/zscalers-ai-bookings-jump-50-130400163.html>